

A/B Testing and Beyond: Designed Experiments for Data Scientists

A Continuing Education Certificate

at

The University of San Francisco's Data Institute

September 4 - October 16, 2018

Instructor: Nathaniel T. Stevens

ntstevens@usfca.edu

Preface

Over the last few decades, there has been an explosion in the amount of data that companies are using to inform decisions. Much of the insight drawn from this influx of data is correlational. Indeed data science is often associated with machine learning, which is powerful in its ability to find patterns and relationships in data for purposes of prediction and classification. However, the ease with which data can be collected provides an enormous opportunity to identify and quantify causal relationships, obtained via experimentation. When causal inference is required, a carefully designed experiment is necessary to evaluate the impact of altering one or more variables on some outcome of interest.

Designed experiments are key to the Scientific Method and are necessary for understanding the world around us. Historically experiments have been used in fields such agriculture, biology, physics, chemistry, pharmacology, epidemiology and industrial engineering, to name a few. More recently however, the utility of designed experiments has been recognized in the world of business and marketing as a tool to increase conversion, strengthen customer retention and improve the bottom line. Companies like Google, Amazon, Facebook, Netflix, Airbnb and Lyft have all adopted experimentation and A/B testing for these purposes. As such, data science practitioners and professionals are beginning to acknowledge experimentation as a foundational tenet of the field.

In this course participants will be exposed to the value experimentation; a strong emphasis is placed on the importance of thinking critically and carefully about the manner in which metrics should be selected and measured, and how data should be collected and analyzed in order to address and answer questions of interest. In particular, this course provides a thorough treatment of available methods and best practices in the design and analysis of experiments. Broad topics include A/B/n testing in which two or more variants are compared,

multivariate experiments such as factorial and fractional-factorial designs, and optimization techniques such as multi-armed bandit experiments and response surface methodology.

What this course does not emphasize is third party experimentation platforms such as Optimizely, Google Analytics, Wasabi, Mixpanel, Apptimize, Split or AB Tasty. While the physical construction of variants and the collection of data is a necessary part of experimentation, there is no standard platform used by all data scientists at all companies. For this reason it would be a poor use of time to train participants in the use of any one platform in particular. The reality is that data scientists will use the experimentation platforms and data pipelines espoused by their own companies.

What this course does emphasize is the statistical principles and practical considerations that underlie effective experimentation. Specifically, participants will develop an appreciation for the careful navigation of the choices and nuances associated with the design of an experiment. Participants will also develop a mastery of the relevant hypothesis tests, power analyses, sample size calculations and analysis methods necessary to draw conclusions and make impactful statements about the question of interest. Participants will also become familiar with using `R` and `Python` to automate components of both the design and the analysis of experiments.

Introduction

In this chapter we discuss what an experiment is, how it differs from other data collection strategies, and why it is so useful. We will also discuss important concepts and important decisions that must be considered when planning an experiment, and we package all of this within a general framework for solving problems and answering questions with planned investigations. First, however, we will lay a foundation of notation and nomenclature which will help to make discussions in this course clear and concise.

1.1 Notation and Nomenclature

In all planned investigations interest lies in solving a problem or answering a particular question using data. The data available for such a task are typically composed of measurements on one or more variables. Here we make a distinction between two classes of variables, based on our interest in them.

The problem/question we wish to address is typically defined in the context of optimizing some metric of interest. In practice such metrics tend to be performance metrics or key performance indicators (KPIs) such as conversion rates, average purchase size, bounce rate, maximum page load time or average session duration, to name just a few. The variable whose measurements are used to calculate such a metric is referred to as the **response variable**. For example, an experiment may involve comparing different messages on a call-to-action button to find which message maximizes the click through rate (CTR). The metric of interest here is the CTR and the corresponding response variable is a binary indicator which identifies whether or not users click the button. As a second example, an experiment may involve the comparison of different webpage designs to decide which one maximizes the aver-

age time on page (TOP). Here the metric of interest is average TOP and the corresponding response variable is the continuous measurement of time on page for each user. Regardless of the type or goal of the experiment, the response variable is the one we are primarily interested in. Throughout this course we will use the letter y to denote response variables.

The variable(s) we believe may influence the response variable are called **explanatory variables** and we tend to think of them as having secondary importance relative to the response variable. In a sense, these are independent variables whereas the response is a dependent variable. In the context of experimentation we refer to explanatory variables as **factors** and we denote them with the letter x . In the simple examples above, the button's message and the webpage's design are the factors that influence CTR and average TOP, respectively.

The different values that a factor takes on in an experiment are referred to as **levels**. Suppose in the button message experiment the following three messages are being tested: “*Submit*”, “*Go*”, and “*Let's Go!*”. In this case the factor ‘button message’ has three levels: $\{Submit, Go, Let's Go!\}$. In the webpage design experiment, suppose two designs are being considered: one with a static image and one with a rotating carousel of static images. In this case the factor ‘webpage design’ has two levels: $\{photo, carousel\}$. It is plain to see that factor levels are what define different **experimental conditions**.

In general, the purpose of an experiment is to alter the levels of one or more factors, and then observe and quantify the resultant effect on the response variable. In order to do this, we must expose **experimental units** to different levels of the factor(s) under study (i.e., to different conditions) and measure their corresponding response value. In the context of online experiments like the examples above, the units are typically users or customers. Suppose that the button in the button message experiment must be clicked in order to complete a digital survey. The users that are exposed to the three different ‘button message’ conditions are the experimental units.

We note briefly that an experiment is not the only way to learn about the relationship between a response variable and one or more factors. In the next section we consider two different data collection strategies and discuss the advantages and disadvantages of each with

respect to understanding the relationship between y and one or more x 's.

1.2 Experiments versus Observational Studies

An **experiment** is composed of a collection of conditions defined by purposeful changes to one or more factors. The goal is to identify and quantify the differences in response variable values across conditions. In other words, the goal is to evaluate the change in response elicited by a change in the factors. In determining whether a factor significantly influences a response, like whether a button's message significantly influences CTR, it is necessary to understand how experimental units respond when exposed to each of the corresponding conditions. However, we cannot simultaneously expose the *same* set of units to each condition; a group of units can be exposed to just one condition. Unfortunately, then, we do not observe how the units respond in the conditions to which they were not exposed. Their hypothetical and unobservable response in these conditions is what we call a **counterfactual**. Because counterfactual outcomes cannot be observed, we require a proxy. Thus, instead, we randomly assign a *different* set of units to each condition and we compare the response variable measurements across conditions. When the units are assigned to the conditions at random, it is reasonable to believe that the only difference between the units in each condition is the fact that they are in different conditions. Thus, if there is a marked difference in the response between the conditions, then this difference can be attributed to the conditions themselves. In this way, we conclude that the observed difference in response values was **caused** by the condition the units were in, and hence by the controlled changes that were made to the factors. The key here is that the factors are purposefully controlled in order to observe the resulting effect on the response.

As mentioned above, generally speaking, the goal in these sorts of investigations is to evaluate the change in response associated with a change in the factors. Strictly speaking one does not require an experiment to do this. Establishing these sorts of relationships can also be done with **observational studies**. The distinction between this and an experiment is that in an observational study there is no measure of control in the data collection process. Instead, data are recorded passively and any relationship between the response and factors

is observed organically. While such an approach provides information about the association between these factors, it does not provide clear information about a causal relationship. When **causal inference** (establishing causal connections between variables) is of interest, it is best if the data arise as a result of an experiment. While methods for establishing causal relationships from observational data do exist (see e.g., propensity score matching ([Rosenbaum and Rubin, 1983](#))), they are much less sound and much more error prone than a carefully designed experiment.

Thus, experiments are advantageous because causal inference is easier than in the context of an observational study. However, experiments can be risky and costly. Consider the situation in which an experimental condition very negatively effects the user experience and results in a revenue loss. This is an outcome, that if at all possible, one would like to avoid.

Another drawback to experimentation is that some experimental conditions may not be ethical. For example, in evaluating whether smoking causes lung cancer, it would be unethical to have a ‘*smoking*’ condition in which subjects are forced to smoke. As a second example, in a pricing experiment it may be perceived as unethical to randomize users to different pricing conditions in which some users pay more money for the same product than others. [Shmueli \(2017\)](#) discusses ethics in online experimentation and points to a recent and controversial emotional contagion experiment at Facebook as being unethical.

While observational studies do not facilitate causal inference as easily as experiments do, they enjoy protection from these other issues since nothing is being manipulated or controlled. Users behave as they normally would and are not forced to participate in something which may be costly or which may be unethical. Thus there is a trade-off between experiments and observational studies: experiments facilitate causal inference, but they can be costly and unethical whereas observational studies are the exact opposite. Thus a data scientist planning an investigation should consider the goals of the investigation and choose their data collection strategy carefully.

In the next section we discuss a framework for planning investigations that formalizes the process by which data is collected to answer questions, regardless of the data collection strategy.

1.3 QPDAC: A Strategy for Answering Questions with Data

In this section we discuss a framework for planning and executing an investigation whose results are in turn analyzed so that conclusions may be drawn about some question of interest. This framework is referred to as QPDAC, an acronym that stands for *Question, Plan, Data, Analysis* and *Conclusion* (Steiner and MacKay, 2005). While this approach is suitable for any formal data-driven investigation, here we emphasize its utility in designing and analyzing experiments. We describe each step of this framework in turn.

Question: Develop a clear statement of the question that needs to be answered. This statement will correspond to some hypothesis that you would like to prove or disprove with an experiment. For example, in the webpage design experiment a question statement might look as follows: “*Relative to the original webpage design with a static image, does a rotating carousel of images decrease bounce rate?*”. It is important that this statement is clear, concise and quantifiable because it will influence many decisions associated with the design and analysis of the experiment. It is also important that everyone involved in the experiment - from data scientists and analysts to product managers and engineers - is aware of the question of interest and hence the goal of the experiment. Experiments may have many goals including, for example, factor screening, optimization or confirmation (we will elaborate on each of these types of experiments as the course progresses). But no matter the goal, it is important that everyone involved is aware of it, and committed to the success of the experiment. Siroker and Koomen (2013) stress the importance of building a culture of testing and experimentation within your organization. When such a culture exists, experimentation is highly valued and can become maximally beneficial. Clearly communicating the question is an excellent first step toward this end.

Plan: In this stage the experiment is designed and all pre-experimental questions should be answered. For example, it is at this stage that the response variable and experimental factors must be chosen. This may seem trivial, but it is arguably the most important step in any experiment and careful consideration should be given to these choices. When choosing the response variable it is important to consider the **Question**; it is through measurements of this variable that the question is answered and so it is necessary to choose a metric that

is related to this question and whose variation can be quantified.

The choice of which factor(s) to manipulate in the experiment will also be guided by the **Question**. Recall that factors are the variables we expect to influence the response. It is important at this stage to brainstorm all such factors that might influence the response and make decisions about whether and how they will be controlled in the experiment. We classify factors into one of three types:

- i. **Design factors:** factors that we will manipulate in the experiment and that define the experimental conditions
- ii. **Nuisance factors:** factors that we expect to influence the response, but whose effect we do not care about. These factors are typically held fixed during the experiment so as to eliminate them as a source of variation in the response variable.
- iii. **Allowed-to-vary factors:** factors that we *cannot* control and factors that we are unaware of. In either case these factors are ones that we do not control in the experiment.

Once these choices have been made it is necessary to define the experimental conditions by deciding which levels of the design factor(s) you will experiment with.

Related to the choice of response variable and design factors is the choice of experimental units. After all, it is the units that are exposed to the different conditions and on which the response variable is measured. In many situations this will be an obvious choice, like an app's users or a company's customers. However, in other situations this decision is not so straightforward. For example, consider online marketplaces like Ebay, Etsy or Airbnb in which it is conceivable that the experimental unit could be the seller/owner or the buyer/renter. The type of question being posed and the particular response variable being measured will typically influence this choice.

With the units defined, conditions established, and the response variable chosen, the final decisions to be made concern the number of units to assign to each condition, and the manner in which this assignment is made. Power analyses and sample size calculations are used to address the former concern and the sampling mechanism addresses the latter. While

random assignment is the standard approach, other hierarchical assignment strategies such as stratified or segmented sampling are also common. We elaborate on these topics later on in the course.

Data: In this stage the data are collected according to the **Plan**. It is extremely important that this step be done correctly; the suitability and effectiveness of the analysis relies on the data being collected correctly. Computer scientists often use the phrase “garbage in, garbage out” to describe the phenomenon whereby poor quality input will always produce faulty output. This sentiment is true here also. If the data quality is compromised, the resulting analysis may be invalid in which case any conclusions drawn will be irrelevant.

One particularly important data quality check is to ensure the assignment strategy is working properly. If the **Plan** requires that units be randomly assigned to conditions, it is prudent to confirm whether condition assignment does appear to be random. A common approach for this is an A/A test, where units are assigned to one of two *identical* conditions. If the assignment was truly random, characteristics of the two groups of units (i.e., measurements of the response variable or demographic composition) should be indistinguishable. If they aren’t, then there is likely something wrong with the assignment mechanism or the manner in which the data are being recorded. Either way, there is a problem that needs to be fixed prior to running the actual experiment.

Analysis: In this stage the **Data** are statistically analyzed to provide an objective answer to the **Question**. This is most typically achieved by way of estimating parameters, fitting models, and carrying out statistical hypothesis tests. If the experiment was well-designed and the data were collected correctly, this step should be straightforward. Throughout the course we will discuss, at length, a variety of statistical analyses whose suitability will depend on the design of the experiment and the type of data that were collected.

Conclusion: In this stage the results of the **Analysis** are considered and one must draw conclusions about what has been learned. These conclusions should then be clearly communicated to all parties involved in - or impacted by - the experiment. Clearly communicating your “wins” or what you learned from your “losses” will help to foster the culture of experimentation [Siroker and Koomen \(2013\)](#) suggest organizations should strive for.

It is very common that these results will precipitate new questions and new hypotheses that further experimentation can help answer. As we will emphasize routinely throughout the course, effective experimentation is sequential; information learned from one experiment helps to inform future experiments and knowledge is generated through a sequence of planned investigations. In this way, the QPDAC framework can be viewed as an ongoing cycle of knowledge generation as illustrated in Figure 1.



Figure 1: QPDAC Cycle

1.4 Fundamental Principles of Experimental Design

Having now described the merits and utility of experimentation, and having provided a framework for planning and executing such an investigation, we now describe three fundamental experimental design principles that should be considered when planning any experiment: *randomization*, *replication*, and *blocking* (Montgomery, 2017). You will see that we have briefly mentioned these concepts previously, but we formalize them here.

Randomization refers both to the manner in which experimental units are selected for inclusion in the experiment and the manner in which they are assigned to experimental

conditions. Note that to avoid the risk of underperforming conditions or conditions with negative side effects, online experiments typically do not include all possible units (users). Instead, some fraction of them is selected for inclusion in the study. Then, once selected, the experimental units are assigned to one of the experimental conditions. Thus we have two levels of randomization.

As we will see later in the course, the validity of many methods of statistical analysis and statistical inference rely on the assumption that inclusion and assignment were done at random. However, there is a more intuitively appealing justification for randomization. The first level of randomization exists to ensure the sample of units included in the experiment is representative of those that were not. This way, the conclusions drawn from the experiment can be generalized to the broader population. The second level of randomization exists to balance out the effects of extraneous variables not under study (i.e., the allowed-to-vary factors). This balancing, in theory, ensures that the units in each condition are as similar to one another as can be, and thus any observed difference in response values can be attributed to the differences between the conditions themselves.

Replication refers to the existence of multiple response observations within each experimental condition and thus corresponds to the situation in which more than one unit is assigned to each condition. Assigning multiple units to each condition provides assurance that the observed results are genuine, and not just due to chance. And as the number of units in each condition increases (i.e., with more replication), we become increasingly sure of the results we observe. For instance, consider the button message experiment introduced previously. Suppose the CTRs in the *Submit*, *Go* and *Let's Go!* conditions were respectively 0.5, 0.5 and 1. If these click-through-rates were calculated from 2 users in each condition, the results would not be nearly as convincing as if they had been calculated from 1000 users in each condition.

The importance of replication likely seems obvious, but the answer to the question “*how much replication is needed?*” is likely less obvious and is just as important. More directly, this question is equivalent to asking “*how many units should be assigned to each condition?*”. The **sample size** for a given condition, denoted by n , is defined to be the number of units

exposed to that condition. We use power analyses and sample size calculations to determine how many units to include in the study, and hence how many response variable observations are necessary to be sufficiently confident in your results. In the context of online experiments, where website traffic is heavy and predictable, replication is often communicated in terms of time as opposed to number of units. For instance, a common question is *“how long does the experiment need to run for?”*. Intuitively, the more confident one wishes to be in the experiment’s results, the larger the sample size needs to be and hence the longer the duration of the experiment. We will formalize these reflections in the chapters to come.

Blocking is the mechanism by which nuisance factors are controlled for. Recall that nuisance factors are known to influence the response variable, but we are not interested in these relationships. Because we wish to ensure the only source of variation in response values is due to the experimental conditions (i.e., changing levels of design factors), we must hold the nuisance factors fixed during the experiment so that they do not impart any variation. Thus we run the experiment at fixed levels of the nuisance factors, i.e., within **blocks**.

For example, consider an email promotion experiment in which the primary goal is to test different variations of the message in the subject line with the goal of maximizing ‘open rate’. However, suppose that it is known that ‘open rate’ is also influenced by the time of day and the day of the week that the email is sent. So as not to conflate the influence of the email’s subject with these time effects, we may elect to send all of the emails at the same time of day and on the same day of the week. Here the block is the particular day and time of day in which the emails are sent. Blocking in this way eliminates these additional sources of variation, and guarantees that observed variation in the response variable is not due to time-of-day or day-of-week effects.

1.5 Exercise: The Instagram Experiment

We end this chapter by pretending we are data scientists at Instagram that need to design an experiment concerning sponsored ads. While ads serve as a source of revenue for Instagram, they also serve as a source of frustration and annoyance to users. Thus,

we would like to run an experiment to gain insight into the interplay between ad revenue, user engagement and factors such as ad frequency, ad type (photo/video), whether the ad's content is targeted or not, etc. Ultimately the goal is to identify a condition that maximizes ad revenue without simultaneously plummeting user engagement below some minimally acceptable threshold.

How would you design such an experiment?

Experiments With Two Conditions

We now consider the design and analysis of an experiment consisting of two experimental conditions – or what many data scientists broadly refer to as “A/B Testing”. Typically the goal of such an experiment is to decide which condition is optimal with respect to some metric of interest. For instance, the canonical A/B test is one in which two versions of a webpage are tested – one with a red button and the other with a blue button – and the ‘winning’ webpage is the one with the button that is clicked most frequently. Although this example is trivial, and it oversimplifies the difficulties, nuances, and importance of such an experiment, it serves as a tangible example of the question being answered: given two options, which one is best?

Formally, such a question is phrased as a statistical hypothesis that we test using the data collected from the experiment. In order to do so we must first define the two experimental conditions by selecting a single design factor and choosing two levels to experiment with. Once these choices are made, the experimental conditions are established and we must randomize n_1 experimental units to one condition and n_2 to the other condition. Next we measure the response variable (y) on each of these units and summarize these response measurements with some metric of interest, θ . Statistically speaking this metric might be a mean, a proportion, a variance, a percentile, or technically any statistic that can be calculated from sample data. Practically speaking such metrics might be things like average-time-on page, average-number-of-bookings, average-number-of-impressions, average-purchase-size, click-through-rate, bounce-rate, conversion-rate, retention-rate, etc. The exact metric chosen will depend on the question being answered and the type of data being collected.

Supposing the metric of interest has been chosen, interest lies in comparing this metric

between the conditions and identifying the optimal condition as the one that optimizes (i.e., maximizes or minimizes) it. Because such a metric is calculated from sample data, which are drawn from a broader population, we view it as an estimate of the corresponding parameter in that population. For example, suppose in the red vs. blue button experiment the click-through-rates of the two conditions are 0.12 (red) and 0.03 (blue). These values are simply sample estimates of the true red vs. blue click-through-rates, which we denote by θ_1 and θ_2 . Thus, $\hat{\theta}_1 = 0.12$ and $\hat{\theta}_2 = 0.03$. Although it is clear that $\hat{\theta}_1 > \hat{\theta}_2$ we must formally decide whether this sample data provides enough evidence to believe that regardless of the sample you might have drawn, the red button is superior to the blue one. In other words, that $\theta_1 > \theta_2$. As mentioned, such a statement is formally phrased as a statistical hypothesis of the form

$$H_0: \theta_1 \leq \theta_2 \text{ vs. } H_A: \theta_1 > \theta_2 \quad (2.1)$$

or

$$H_0: \theta_1 \geq \theta_2 \text{ vs. } H_A: \theta_1 < \theta_2 \quad (2.2)$$

Since it is the null hypothesis H_0 that is assumed to be true at baseline, which statement one wishes to test depends on this baseline assumption. However, notice that H_0 and H_A are complements of one another, and so only one of them is true. Furthermore, our decision is also binary: based on the observed data we choose to reject or not reject H_0 . Thus, regardless of which direction you choose to state your hypothesis, the conclusion you draw will be the same. To make this clear, suppose that in the red vs. blue button experiment θ_1 and θ_2 respectively represent the click-through-rates for the red and blue buttons. If the data suggest red is better, it doesn't matter which hypothesis statement we test. If we test (2.1) we will reject H_0 (and conclude that red is best), and if we test (2.2) we will not reject H_0 (and hence conclude that red is best).

Note that when comparing some metric of interest across two conditions, it may also be of interest to test the two-sided hypothesis

$$H_0: \theta_1 = \theta_2 \text{ vs. } H_A: \theta_1 \neq \theta_2. \quad (2.3)$$

This hypothesis provides no information about which of the two conditions is best, but does tell us whether they are different. As such, it may be used as an initial check of whether the conditions are different at all. If they aren't, then there is no reason to proceed. But if they are, then we would use hypothesis (2.1) or (2.2) to help determine the optimal condition. For a general review of statistical inference and hypothesis testing, please refer to Appendix A.2.

In the context of hypotheses such as (2.1), (2.2) and (2.3), we discuss in this chapter how to design an experiment to test them and we discuss how to analyze observed data to formally draw conclusions about them. In particular we discuss how to choose the number of units to assign to each condition, and we describe a variety of analysis techniques appropriate for different metrics of interest, and different types of response variables.

2.1 Comparing Means in Two Conditions

In this section we restrict attention to the situation in which the response variable of interest is measured on a continuous scale, although the associated methodology is also commonly applied when response variables are discrete and, for example, represent counts (as in the number of times an event of interest occurs). In these cases we assume that the response observations collected in the two conditions follow normal distributions, and in particular

$$Y_{i1} \sim N(\mu_1, \sigma^2) \text{ and } Y_{i2} \sim N(\mu_2, \sigma^2)$$

where $i = 1, 2, \dots, n_j$ for $j = 1, 2$. Thus Y_{ij} represents the response observation for the i^{th} unit in the j^{th} condition, and we assume that the measurements in the two conditions could reasonably have been drawn from a normal distribution with mean μ_1 (in the first condition) or μ_2 (in the second) and common variance σ^2 . Thus we believe that the distributions from which these samples were drawn only differ (if they differ at all) with respect to the mean, and in no other way. Thus a comparison of the two conditions corresponds to a comparison of the expected responses (i.e., the means) in each of them. Specifically we test hypotheses

of the form

$$H_0: \mu_1 = \mu_2 \text{ vs. } H_A: \mu_1 \neq \mu_2 \quad (2.4)$$

$$H_0: \mu_1 \leq \mu_2 \text{ vs. } H_A: \mu_1 > \mu_2 \quad (2.5)$$

$$H_0: \mu_1 \geq \mu_2 \text{ vs. } H_A: \mu_1 < \mu_2 \quad (2.6)$$

In the following subsections we describe how to analyze data of this form and draw conclusions about such hypotheses and we also describe how to choose the sample size that allows one to be sufficiently confident in their conclusions.

2.1.1 The Two-Sample t -Test

In order to test hypotheses (2.4), (2.5) and (2.6) we must first calculate a **test statistic**. Because $Y_{ij} \sim N(\mu_j, \sigma^2)$, it is also true that $\bar{Y}_j = \frac{1}{n_j} \sum_{i=1}^{n_j} Y_{ij} \sim N(\mu_j, \frac{\sigma^2}{n_j})$ and hence that

$$\frac{(\bar{Y}_1 - \bar{Y}_2) - (\mu_1 - \mu_2)}{\sigma \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}} \sim N(0, 1). \quad (2.7)$$

Although we can substitute a hypothesized value for $\mu_1 - \mu_2$ into this expression, we do not have a hypothesized value for σ . As such, we replace it in the equation above using the following estimate

$$\hat{\sigma}^2 = \frac{\sum_{i=1}^{n_1} (Y_{i1} - \bar{Y}_1)^2 + \sum_{i=1}^{n_2} (Y_{i2} - \bar{Y}_2)^2}{n_1 + n_2 - 2}.$$

Note that this quantity is simply a pooled estimate of σ^2 based on the sample variances in the two conditions.

Substituting $\hat{\sigma}$ for σ gives

$$T = \frac{(\bar{Y}_1 - \bar{Y}_2) - (\mu_1 - \mu_2)}{\hat{\sigma} \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}} \sim t_{(n_1+n_2-2)} \quad (2.8)$$

which is the test statistic for these hypothesis tests and where $t_{(n_1+n_2-2)}$ is the **null distribution**. It is for this reason that the test is called a “*t*-test”.

Hypotheses (2.4), (2.5) and (2.6) are formally tested by calculating the observed value of T from our sample data $\{y_{11}, y_{21}, \dots, y_{n_11}\}$ and $\{y_{12}, y_{22}, \dots, y_{n_22}\}$ and evaluating its extremity in the context of the $t_{(n_1+n_2-2)}$ distribution. Given the sample data, we have $\bar{y}_1 = \hat{\mu}_1$ and $\bar{y}_2 = \hat{\mu}_2$ and so the observed test statistic is given by

$$\begin{aligned} t &= \frac{(\bar{y}_1 - \bar{y}_2) - (\mu_1 - \mu_2)}{\hat{\sigma} \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}} \\ &= \frac{(\hat{\mu}_1 - \hat{\mu}_2) - 0}{\hat{\sigma} \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}} \\ &= \frac{(\hat{\mu}_1 - \hat{\mu}_2)}{\hat{\sigma} \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}}. \end{aligned} \quad (2.9)$$

Notice that we have also substituted the hypothesized value $\mu_1 - \mu_2 = 0$ indicating a null assumption of ‘no difference’ between the two conditions.

To decide whether to reject H_0 or not reject H_0 we must calculate the **p-value** – the probability of observing a value of the test statistic at least as extreme as the one we observed, if H_0 really were true. In the case of hypothesis (2.4) the p-value is calculated as p-value = $2P(T \geq |t|)$; in the case of hypothesis (2.5) the p-value is calculated as p-value = $P(T \geq t)$, and in the case of hypothesis (2.6) the p-value is calculated as p-value = $P(T \leq t)$, where here $T \sim t_{(n_1+n_2-2)}$. See Figure A.9 for a visual depiction of these calculations. We then decide to reject (or not reject) H_0 on the basis of a comparison between the calculated p-value and the **significance level** α . If p-value $\leq \alpha$ we reject H_0 in favor of H_A , and if p-value $> \alpha$ we do not reject H_0 .

Next we consider an example to illustrate these ideas more concretely.

2.1.2 Example: Instagram Ad Frequency

Suppose, again, that you are a data scientist at Instagram, and you are interested in running an experiment to learn about how user engagement is influenced by ad frequency. Currently users see one ad every 8 posts in their social feed, but, in order to increase ad revenue, your manager is pressuring your team to show one ad every 5 posts, under the assumption that users will not behave any differently under this new regime. You are justifiably nervous about this change and you worry that this will substantially decrease user engagement and hurt the overall user experience. As such you propose an experiment to test this new regime before rolling it out to all users. The experiment you propose is an A/B test in which average session time (i.e., the length of time a user engages with the app – in minutes) is compared between the two ad frequency conditions. You hypothesize that the current ad frequency (condition 1) will correspond to a significantly longer average session time than the proposed ad frequency (condition 2).

Thus, in the language and notation of these notes, you're interested in testing a hypothesis such as (2.5) where μ_1 represents the average session time of a user in the 7:1 ad frequency condition and μ_2 represents the average session time of a user in the 4:1 ad frequency condition. The null hypothesis here assumes what your manager assumes – that increased ad frequency does not lead to reduced engagement ($H_0: \mu_1 \leq \mu_2$). Thus you expect to collect data that contradicts this statement so that it can be rejected in favor of the alternative that says that increased ad frequency significantly reduces the amount of time users are engaged with the app ($H_A: \mu_1 > \mu_2$).

In order to test this hypothesis you randomize $n_1 = 500$ users to the 7:1 ad frequency condition and $n_2 = 500$ users to the 4:1 condition. The data you collect is summarized as follows: The average session time in the 4:1 condition is $\hat{\mu}_1 = \bar{y}_1 = 4.9162$ with a standard deviation of $s_1 = 0.9634$, and in the 7:1 condition the average session time is $\hat{\mu}_2 = \bar{y}_2 = 3.0518$

with a standard deviation of $s_2 = 0.9950$. The pooled standard deviation estimate is

$$\hat{\sigma} = \sqrt{\frac{499 \cdot 0.9634^2 + 499 \cdot 0.995^2}{998}} = 0.9793.$$

These summaries support your suspicion: session time appears to be negatively effected by an increased ad frequency.

To determine whether this difference is statistically significant, you formally test the hypothesis by calculating a p-value. To do this, you must first calculate the observed test statistic. Substituting these summaries into equation (2.9) gives

$$t = \frac{4.9162 - 3.0518}{0.9793 \sqrt{\frac{2}{500}}} = 30.1013.$$

The p-value associated with this test is $P(T \geq 30.1013)$ where $T \sim t_{998}$. When calculated this probability is equal to 1.84×10^{-142} , which is essentially 0. In R this probability is calculated using the command `pt(30.1013, df = 998, lower.tail = F)`. We can also use the `t.test()` function in R to do the whole test; you need only pass it the data and a few other arguments and it will calculate the necessary summaries, the test statistic and the p-value. Note that to replicate the results here we must set the logical argument `var.equal` to `TRUE`. We discuss an alternative approach to take when the variances are not assumed to be equal in Section 2.1.4.

In order to draw a conclusion, we must compare our calculated p-value to the significance level $\alpha = 0.05$. Since $1.84 \times 10^{-142} < 0.05$ we reject the null hypothesis in favor of the alternative. In the context of the experiment, this means that increased ad frequency significantly reduces the amount of time users engage with the app. In particular, you can expect a 1 minute and 52 second reduction in average session time when you move from a 4:1 ad frequency to a 7:1 frequency.

In fact, depending on the speed a user scrolls through their feed, this this increased ad frequency could actually reduce ad revenue; suppose that the typical user spends roughly 5 seconds looking at each post, which means they scroll through 12 posts per minute. In the

7:1 ad frequency condition a user would then see 1.5 ads per minute, and in the 4:1 frequency condition a user would see 2.4 ads per minute. Although a user in the 7:1 condition sees fewer ads per minute, they spend more time on the app. At an average session time of roughly 5 minutes, they see 7.5 ads per session, whereas a user in the 4:1 condition, whose session duration is roughly 3 minutes, will see 7.2 ads per minute. As such, it would be ill-advised to adopt this new ad regime from both the perspective of user engagement and ad revenue.

2.1.3 Power Analysis and Sample Size Calculations

When designing a two-condition experiment (i.e., an A/B test), the most important question (once the response variable and conditions have been chosen) is “*How many units do I need in each condition?*”. The answer to this question is determined by the frequency with which we are comfortable drawing the wrong conclusion.

Recall that because H_0 and H_A are complements of one another, exactly one of them is correct. Thus, when we choose to reject or not reject H_0 we risk drawing the wrong conclusion. In this context we can make two types of errors:

- Type I Error: Reject H_0 when it is in fact true
- Type II Error: Do not reject H_0 when it is in fact false

Ideally these types of errors would happen very infrequently. Fortunately we are able to control the frequency with which such errors are made through the **significance level** and the **power** of the hypothesis test. The significance level is denoted by α where

$$\alpha = P(\text{Type I Error}) = P(\text{Reject } H_0 | H_0 \text{ is true})$$

and the power of the test is denoted by $1 - \beta$ where

$$\beta = P(\text{Type II Error}) = P(\text{Do not reject } H_0 | H_0 \text{ is false}).$$

Thus, a test that has a small significance level and large power is desirable as it simultaneously

minimizes the chances of committing both Type I and Type II errors.

In practice these values are chosen to be consistent with one's risk tolerance, though $\alpha = 0.05$ and $\beta = 0.2$ are the standard choices. As we will show here, the significance level and power of a hypothesis test are related to one another and also related to the sample size. In fact, for a given sample size, as the chances of a Type I error decrease, the chances of a Type II error increase, and vice versa. However, if we want to fix α and β at particular values, we can determine what sample size is necessary to do so. Thus, it is important to understand the nature of the relationship between these quantities – if you alter one of them, the others will also change.

In what follows, we will derive the formula that quantifies this relationship, and which can be used to determine the sample size necessary to keep Type I and Type II errors at bay. We do so assuming that the parameter σ is known, or at least that we have a reasonable guess as to what it might be. Note that we do not need to make this assumption once the data are collected, but we do need to make it prior to data collection. Thus for the development below we define our test statistic T as in equation (2.7) which means that we will be working with the $N(0, 1)$ distribution.

We begin by precisely defining what it means (in terms of the test statistic) to reject H_0 . In all cases this happens when $p\text{-value} \leq \alpha$. In the context of a two-sided hypothesis such as (2.4) this happens when $t \geq z_{\alpha/2}$ or $t \leq -z_{\alpha/2}$, where $z_{\alpha/2}$ is the $(1 - \alpha/2)^{th}$ quantile of the standard normal distribution. Thus we can define a **rejection region** $R = \{t \mid t \geq z_{\alpha/2} \text{ or } t \leq -z_{\alpha/2}\}$ that describes all values of t for which H_0 would be rejected. Similar rejection regions can be defined for hypotheses (2.5) and (2.6) as well. These are respectively given by $R = \{t \mid t \geq z_{\alpha/2}\}$ and $R = \{t \mid t \leq -z_{\alpha/2}\}$. All of these rejection regions are depicted in blue in Figure 2.

Having defined these we now derive the formula which, for a given significance level and power, prescribes how many units should be assigned to each condition. Although it is very common to assign the same number of units to each of the conditions (i.e., $n_1 = n_2$), we will keep this derivation general and not make this specific requirement. What we do require, however, is an assumption about the relative sizes of n_1 and n_2 . Specifically, we need to




Figure 2: Rejection regions corresponding to one and two-sided hypotheses

specify k where $n_1 = kn_2$. In the case that equal sample sizes are desired we would simply take $k = 1$. Furthermore, we provide this derivation under the assumption that we are dealing with a hypothesis that has a two-sided alternative as in (2.4). We indicate where and how this derivation changes if sample size calculations in the context of a one-sided hypothesis test is of interest.

We begin by considering the power of the hypothesis test:

$$\begin{aligned}
1 - \beta &= P(\text{Reject } H_0 \mid H_0 \text{ is false}) \\
&= P(T \in R \mid H_0 \text{ is false}) \text{ where } R \text{ is the rejection region} \\
&= P(T \geq z_{\alpha/2} \text{ or } T \leq -z_{\alpha/2} \mid H_0 \text{ is false}) \\
&= P(T \geq z_{\alpha/2} \mid H_0 \text{ is false}) + P(T \leq -z_{\alpha/2} \mid H_0 \text{ is false}) \\
&= P\left(\frac{(\bar{Y}_1 - \bar{Y}_2)}{\sigma\sqrt{\frac{1}{n_1} + \frac{1}{n_2}}} \geq z_{\alpha/2} \mid H_0 \text{ is false}\right) + P\left(\frac{(\bar{Y}_1 - \bar{Y}_2)}{\sigma\sqrt{\frac{1}{n_1} + \frac{1}{n_2}}} \leq -z_{\alpha/2} \mid H_0 \text{ is false}\right)
\end{aligned}$$

If $H_0 : \mu_1 = \mu_2$ were true, and hence $\mu_1 - \mu_2 = 0$ were true, then the ratios in the preceding line would follow a $N(0, 1)$ distribution. However, we know that H_0 is false which means that $\mu_1 - \mu_2 = \delta$ for some none-zero δ , and so it is

$$\frac{(\bar{Y}_1 - \bar{Y}_2) - \delta}{\sigma\sqrt{\frac{1}{n_1} + \frac{1}{n_2}}}$$

that follows a $N(0, 1)$ distribution. Let us make this substitution, being sure to replicate what is done on the left side of inequalities on the right. Also note that we no longer need to write “ $|H_0$ is false” since we are now exploiting this fact.

$$\begin{aligned} 1 - \beta &= P \left(\frac{(\bar{Y}_1 - \bar{Y}_2) - \delta}{\sigma \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}} \geq z_{\alpha/2} - \frac{\delta}{\sigma \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}} \right) + P \left(\frac{(\bar{Y}_1 - \bar{Y}_2) - \delta}{\sigma \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}} \leq -z_{\alpha/2} - \frac{\delta}{\sigma \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}} \right) \\ &= P \left(Z \geq z_{\alpha/2} - \frac{\delta}{\sigma \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}} \right) + P \left(Z \leq -z_{\alpha/2} - \frac{\delta}{\sigma \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}} \right) \text{ where } Z \sim N(0, 1) \end{aligned}$$

Note that depending on the sign of δ , just one of these terms will dominate. To see this, suppose $\delta > 0$; then $-z_{\alpha/2} - \frac{\delta}{\sigma \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}}$ will be an extremely negative number and the probability that a standard normal random variable is smaller than an extremely negative number is effectively 0, and only the first term remains. Now suppose $\delta < 0$; then $z_{\alpha/2} - \frac{\delta}{\sigma \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}}$ will be an extremely positive number and the probability that a standard normal random variable is larger than an extremely positive number is effectively 0, and only the second term remains. Assume, without loss of generality, that $\delta > 0$ in which case

$$1 - \beta = P \left(Z \geq z_{\alpha/2} - \frac{\delta}{\sigma \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}} \right)$$

Because this probability is equal to $1 - \beta$ we know that $z_{\alpha/2} - \frac{\delta}{\sigma \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}}$ must be equal to $z_{1-\beta}$, the β^{th} quantile of the standard normal distribution. Thus

$$z_{1-\beta} = z_{\alpha/2} - \frac{\delta}{\sigma \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}}$$

and we can rearrange this equation solving for the sample size. But first we must substitute $n_1 = kn_2$ so that there is just a single sample size to solve for:

$$z_{1-\beta} = z_{\alpha/2} - \frac{\delta}{\sigma \sqrt{\frac{1}{kn_2} + \frac{1}{n_2}}} = z_{\alpha/2} - \frac{\sqrt{n_2}\delta}{\sigma \sqrt{\frac{1}{k} + 1}}$$

Solving for n_2 yields:

$$n_2 = \frac{(\frac{1}{k} + 1)(z_{\alpha/2} - z_{1-\beta})^2 \sigma^2}{\delta^2} \quad (2.10)$$

and then n_1 is found by computing kn_2 . When equal sample sizes are desired ($k = 1$) each condition receives n units where

$$n = \frac{2(z_{\alpha/2} - z_{1-\beta})^2 \sigma^2}{\delta^2}. \quad (2.11)$$

So when calculating a sample size we need to have chosen α and β (our Type I and Type II error rates), we need a guess as to what σ is, and we need a value for δ . With all of this information one can readily use the formulae above to calculate n_1 and n_2 .

But where does the δ value come from? We define δ to be the **effect size** of the test. The effect size for hypothesis tests like (2.4), (2.5) or (2.6) refers to the minimal difference between conditions (i.e., between μ_1 and μ_2) that we find to be practically relevant and that we would like to detect as being statistically significant. For instance, imagine we are comparing the average length of time users spend engaging with their Instagram apps as in Section 2.1.2. Suppose that condition 1 (7:1 ad frequency) corresponds to the current version of the app, and you know users engage with the app for an average of 5 minutes. Now suppose that condition 2 corresponds to the 4:1 ad frequency. Would it be practically relevant if users in condition 2 spend an average of 4.8 minutes engaged with the app? If not, would it be practically important if these users spent an average of 3.5 minutes engaged with the app? The answer to the question “*What is the minimal difference between μ_1 and μ_2 that is practically important?*” is effect size, and is what is captured by $\delta = \mu_1 - \mu_2$.

Sometimes effect size is defined on a standardized scale, and communicated in numbers of standard deviations as opposed to the absolute scale described above. In this case δ is defined as

$$\delta = \frac{\mu_1 - \mu_2}{\sigma}$$

and the sample size formula (2.10) simplifies to

$$n_2 = \frac{(\frac{1}{k} + 1)(z_{\alpha/2} - z_{1-\beta})^2}{\delta^2}$$

and the sample size formula (2.11) simplifies to

$$n = \frac{2(z_{\alpha/2} - z_{1-\beta})^2}{\delta^2}.$$

The advantage of defining effect size on a standardized scale is that we do not require knowing or guessing σ in our sample size calculations.

It is important to also consider how these formulae change if we were performing sample size calculations for one-sided hypothesis tests. In these cases the rejection regions are also one-sided and based on the quantile z_α instead of $z_{\alpha/2}$. It turns out that this is the only difference, and when carried through the derivation yields sample size formulae equivalent to (2.10) and (2.11) but with $z_{\alpha/2}$ replaced by z_α .

As should be clear by looking at equations (2.10) and (2.11), there is an interdependent relationship between sample size, significance level, power, and effect size. These equations can be rearranged to isolate for any of these variables, which illustrates the fact that changing one variable leads to a change in all of the others. For an interactive demonstration of these interdependencies feel free to tinker with the sample size calculator found at the following link: <https://nathaniel-t-stevens.shinyapps.io/SampleSizeCalculator/>.

2.1.4 When Assumptions are Invalid

When testing hypotheses of the form (2.4), (2.5) and (2.6) using the two-sample t -test described in Section 2.1.1, we make two key assumptions. First, we assume that the variance in the two conditions are equal, and second, we assume that the response observations in each condition follow a normal distribution. In this subsection we describe alternative approaches when these assumptions are not valid. We begin with the equal variance assumption.

Welch’s t -Test: When it is unreasonable to assume that the response variable mea-

measurements in each condition have equal variances, an approach that accommodates $Y_{ij} \sim N(\mu_j, \sigma_j^2)$ for $j = 1, 2$, and hence $\sigma_1^2 \neq \sigma_2^2$, is to be preferred. In this situation we may use the test statistic

$$t = \frac{(\hat{\mu}_1 - \hat{\mu}_2)}{\sqrt{\frac{\hat{\sigma}_1^2}{n_1} + \frac{\hat{\sigma}_2^2}{n_2}}}$$

where $\hat{\sigma}_j^2$ is the sample variance of the response measurements in condition $j = 1, 2$. However, this statistic does not follow a t -distribution exactly; it *approximately* follows a t -distribution with

$$\nu = \frac{\left(\frac{\hat{\sigma}_1^2}{n_1} + \frac{\hat{\sigma}_2^2}{n_2}\right)^2}{\frac{(\hat{\sigma}_1^2/n_1)^2}{n_1-1} + \frac{(\hat{\sigma}_2^2/n_2)^2}{n_2-1}}$$

degrees of freedom. Carrying out the test using a $t_{(\nu)}$ null distribution (with ν as above) is referred to as *Welch's t -test* after Bernard L. Welch who devised this approximation ([Welch, 1947](#)). This test can be carried out in R using the `t.test()` function but with the logical argument `var.equal` set to `FALSE`.

In order to decide whether $\sigma_1^2 \neq \sigma_2^2$ and hence whether Welch's t -test is necessary, one might consider formally testing the hypothesis

$$H_0: \sigma_1^2 = \sigma_2^2 \text{ vs. } H_A: \sigma_1^2 \neq \sigma_2^2 \quad (2.12)$$

Such a hypothesis is commonly tested using an **F -test of equal variances**. The F -test assumes that $Y_{ij} \sim N(\mu_j, \sigma_j^2)$ which consequently means that

$$\frac{(n_j - 1)\hat{\sigma}_j^2}{\sigma_j^2} \sim \chi_{n_j-1}^2$$

and hence that

$$T = \frac{\hat{\sigma}_1^2/\sigma_1^2}{\hat{\sigma}_2^2/\sigma_2^2} \sim F(n_1 - 1, n_2 - 1).$$

Assuming H_0 is true, $\sigma_1^2/\sigma_2^2 = 1$ and so the observed value of the test statistic is

$$t = \frac{\hat{\sigma}_1^2}{\hat{\sigma}_2^2}$$

which we compare to the null distribution $F(n_1 - 1, n_2 - 1)$. Note that it is because the null

distribution is an F -distribution that this test is known as an F -test.

Because the F -distribution is not symmetrical and not defined for negative values, in the context of the two-sided hypothesis above the p-value is calculated to be

$$\text{p-value} = P(T \geq t) + P(T \leq 1/t)$$

since values greater than or equal to t and less than or equal to $1/t$ are what is considered “at least as extreme” in this situation. One-sided alternatives might also be considered where $H_A: \sigma_1^2 > \sigma_2^2$ or $H_A: \sigma_1^2 < \sigma_2^2$ in which case the p-values are respectively defined as $\text{p-value} = P(T \geq t)$ and $\text{p-value} = P(T \leq t)$. This test can be carried out in R using the `var.test()` function.

Permutation and Randomization Tests: All of the previous tests assume the response measurements are normally distributed. However, many situations exist in which a numeric response variable does not follow a normal distribution. Using the observed data, this assumption can be informally evaluated using QQ-plots or histograms, or formally using the Shapiro-Wilk test ([Shapiro and Wilk, 1965](#)). While both the Student’s t -test and Welch’s t -test are fairly robust to non-normality, it would be preferable to have a test that does not rely on this assumption. *Permutation* and *randomization tests* are nonparametric resampling techniques that may be used for this purpose in this context.

Suppose you collect response measurements $\{y_{11}, y_{21}, \dots, y_{n_11}\}$ and $\{y_{12}, y_{22}, \dots, y_{n_22}\}$ in conditions 1 and 2, respectively. Using these measurements you then estimate some metric of interest θ in the two conditions yielding $\hat{\theta}_1$ and $\hat{\theta}_2$. The goal, then, is to compare $\hat{\theta}_1$ to $\hat{\theta}_2$ in accordance with hypotheses such as (2.1), (2.2) or (2.3) to decide whether $\theta_1 = \theta_2$, $\theta_1 > \theta_2$ or $\theta_1 < \theta_2$. The philosophy behind the resampling approaches to testing such hypotheses is described below.

If H_0 is true and there is truly no difference between the conditions, then the samples $\{y_{11}, y_{21}, \dots, y_{n_11}\}$ and $\{y_{12}, y_{22}, \dots, y_{n_22}\}$ should be very similar and permuting the labels ‘condition 1’ and ‘condition 2’ associated with each response measurement should not sub-

stantially change $\hat{\theta}_1$ or $\hat{\theta}_2$. In fact, if the null hypothesis is true, each of the

$$\binom{n_1 + n_2}{n_1} = \binom{n_1 + n_2}{n_2}$$

arrangements of the observed data are equally likely. A true **permutation test** takes the test statistic to be $t = \hat{\theta}_1 - \hat{\theta}_2$ and then takes as the null distribution the set of test statistics calculated on each of the $\binom{n_1 + n_2}{n_1} = \binom{n_1 + n_2}{n_2}$ arrangements of data. A formal conclusion about H_0 is drawn on the basis of the extremity of t in the context of this null distribution. The p-value associated with such a test is calculated empirically as the proportion of resampled test statistics that were “at least as extreme” as t .

While conceptually appealing, the permutation test is not practical in most circumstances because the number of permutations of the data becomes enormous, even for relatively small sample sizes. For instance, if $n_1 = n_2 = 50$, there are $\binom{100}{50} = 1.09 \times 10^{29}$ distinct arrangements of the data. Thus, since true permutation tests tend to be computationally expensive, a practical approximation is the **randomization test** which simply investigates a large number of resamples, as opposed to all possible permutations. An algorithm for performing a randomization tests is as follows:

1. Calculate the test statistic $t = \hat{\theta}_1 - \hat{\theta}_2$ on the original sample.
2. Resample the data without replacement so that n_1 observations are randomly associated with a resampled ‘condition 1’: $\{y_{11}^*, y_{21}^*, \dots, y_{n_{11}}^*\}$ and n_2 observations are randomly associated with a resampled ‘condition 2’: $\{y_{12}^*, y_{22}^*, \dots, y_{n_{22}}^*\}$.
3. Calculate the value of the test statistic, labeled t^* , on this resampled data.
4. Repeat steps 2 and 3 B times ($B = 1000$ or 2000 are common choices).
5. Compare t to the null distribution which is derived from the B resampled values of t^* , and calculate the p-value.

The p-values associated with tests of this sort are calculated differently depending on whether the alternative hypothesis, H_A , is one- or two-sided. These calculations are summarized below:

- $H_A: \theta_1 \neq \theta_2$: p-value = The proportion of resampled test statistics $t^* \geq |t|$ or $\leq -|t|$
- $H_A: \theta_1 > \theta_2$: p-value = The proportion of resampled test statistics $t^* > t$
- $H_A: \theta_1 < \theta_2$: p-value = The proportion of resampled test statistics $t^* < t$

See [Edgington and Onghena \(2007\)](#) for a more thorough and general treatment of randomization tests.

2.2 Comparing Proportions in Two Conditions

Very often the response variable in an A/B test is binary, indicating whether an experimental unit did, or did not, perform some action of interest. In cases like these we let

$$Y_{ij} = \begin{cases} 1, & \text{if unit } i \text{ in condition } j \text{ performs the action of interest} \\ 0, & \text{if unit } i \text{ in condition } j \text{ does not perform the action of interest} \end{cases}$$

for $i = 1, 2, \dots, n_j$, $j = 1, 2$. Examples of “actions of interest” include opening an email, clicking a button, watching an ad, leaving a webpage without interacting with it, etc. In each case unit i ’s response variable is recorded as a 1 if they perform the action and a 0 otherwise. Interest then lies in deciding which condition is optimal, where the optimal condition is the one for which the likelihood that a unit performs the action is highest (when maximization is of interest) or smallest (when minimization is of interest).

To formally decide which condition is optimal we must make an assumption about the distribution of the response variable. Because the Y_{ij} ’s are binary, it is common to assume that they follow a Bernoulli distribution:

$$Y_{ij} \sim \text{BIN}(1, \pi_j)$$

where π_j represents the probability that $Y_{ij} = 1$, i.e., the probability that a unit in condition j performs the action of interest. The goal of the experiment then, is to determine whether

$\pi_1 = \pi_2$, $\pi_1 > \pi_2$ or $\pi_1 < \pi_2$. This decision is formally made in association with the following hypotheses:

$$H_0: \pi_1 = \pi_2 \text{ vs. } H_A: \pi_1 \neq \pi_2 \quad (2.13)$$

$$H_0: \pi_1 \leq \pi_2 \text{ vs. } H_A: \pi_1 > \pi_2 \quad (2.14)$$

$$H_0: \pi_1 \geq \pi_2 \text{ vs. } H_A: \pi_1 < \pi_2 \quad (2.15)$$

In the subsections that follow we describe how to analyze data of this form and draw conclusions about hypotheses like these. We also describe power analyses and sample size calculations in this context as well.

2.2.1 The Z-test for Proportions

In order to test hypotheses (2.13), (2.14) and (2.15) we must calculate a test statistic. Due to the **Central Limit Theorem**¹ we know that for large enough n_j the random variable $\bar{Y}_j = \frac{1}{n_j} \sum_{i=1}^{n_j} Y_{ij} \sim N(\pi_j, \frac{\pi_j(1-\pi_j)}{n_j})$. Thus, with a large amount of replication \bar{Y}_{ij} will approximately follow a normal distribution. Based on this result

$$\frac{(\bar{Y}_1 - \bar{Y}_2) - (\pi_1 - \pi_2)}{\sqrt{\frac{\pi_1(1-\pi_1)}{n_1} + \frac{\pi_2(1-\pi_2)}{n_2}}} \sim N(0, 1). \quad (2.16)$$

As a general rule of thumb, this approximation may be very poor unless $n_j\pi_j \geq 10$ and $n_j(1 - \pi_j) \geq 10$ for both $j = 1, 2$.

Although we can substitute a hypothesized value for $\pi_1 - \pi_2$ (i.e., zero) into the equation above, we have no hypothesized value for π_1 or π_2 individually, and so this equation is not

¹The Central Limit Theorem states that for any sequence of random variables X_1, X_2, \dots, X_n with $E[X_i] = \mu$ and $Var[X_i] = \sigma^2 < \infty$ for each $i = 1, 2, \dots, n$, the random variable \bar{X} follows a $N(\mu, \sigma^2)$ distribution for large enough n (i.e., as $n \rightarrow \infty$).

calculable in practice. As such we replace instances of π_1 and π_2 in the denominator with estimates. Because $\pi_1 = \pi_2 = \pi$ under the null hypothesis, we use the pooled estimate given by

$$\hat{\pi} = \frac{n_1 \hat{\pi}_1 + n_2 \hat{\pi}_2}{n_1 + n_2} \quad (2.17)$$

where $\hat{\pi}_1$ and $\hat{\pi}_2$ are respectively equal to \bar{Y}_1 and \bar{Y}_2 . We note that when the response variable is binary, means equate to proportions and so hypothesis tests in this setting amount to a comparison of proportions.

Making these substitutions gives

$$T = \frac{(\bar{Y}_1 - \bar{Y}_2) - (\pi_1 - \pi_2)}{\sqrt{\hat{\pi}(1 - \hat{\pi}) \left(\frac{1}{n_1} + \frac{1}{n_2} \right)}} \quad (2.18)$$

which also approximately follows a $N(0, 1)$ distribution. Thus T is the test statistic associated with hypotheses (2.13), (2.14) and (2.15) where $N(0, 1)$ is the null distribution. It is for this reason that the test is called a “Z-test”.

To formally test these hypotheses we calculate the observed value of the test statistic, t , from our sample data $\{y_{11}, y_{21}, \dots, y_{n_1 1}\}$ and $\{y_{12}, y_{22}, \dots, y_{n_2 2}\}$ and evaluate its extremity in the context of the $N(0, 1)$ distribution. Given the sample data, we have $\bar{y}_1 = \hat{\pi}_1$ and $\bar{y}_2 = \hat{\pi}_2$ and so the observed test statistic is given by

$$\begin{aligned} t &= \frac{(\bar{y}_1 - \bar{y}_2) - (\pi_1 - \pi_2)}{\sqrt{\hat{\pi}(1 - \hat{\pi}) \left(\frac{1}{n_1} + \frac{1}{n_2} \right)}} \\ &= \frac{(\hat{\pi}_1 - \hat{\pi}_2) - 0}{\sqrt{\hat{\pi}(1 - \hat{\pi}) \left(\frac{1}{n_1} + \frac{1}{n_2} \right)}} \\ &= \frac{(\hat{\pi}_1 - \hat{\pi}_2)}{\sqrt{\hat{\pi}(1 - \hat{\pi}) \left(\frac{1}{n_1} + \frac{1}{n_2} \right)}} \end{aligned} \quad (2.19)$$

Notice that we have also substituted the hypothesized value $\pi_1 - \pi_2 = 0$ indicating a null

assumption of ‘no difference’ between the two conditions.

As is typical, we decide whether to reject or not reject H_0 based on the size of the test’s p-value in relation to the significance level α . If p-value $\leq \alpha$ we reject H_0 in favor of H_A , and if p-value $> \alpha$ we do not reject H_0 . The p-values associated with hypotheses (2.13), (2.14) and (2.15) are respectively calculated as $2P(T \geq |t|)$, $P(T \geq t)$, and $P(T \leq t)$ where in each case $T \sim N(0, 1)$.

In the next subsection we consider an example to illustrate these ideas more concretely.

2.2.2 Example: Optimizing Optimizely

Siroker and Koomen (2013) discuss an A/B test they ran on the Optimizely website. In particular they were in the midst of a complete website redesign, and they were interested in how new versions of certain pages influenced things like conversion and engagement relative to the old version. One such metric they were interested in was whether or not the redesigned homepage lead to a significant increase in the number of new accounts created.

Thus, in the language and notation of these notes, they were interested in testing a hypothesis such as (2.15) where π_1 represents the probability that a user would create an account on the old homepage and π_2 represents the probability that a user would create an account while viewing the redesigned homepage. The null hypothesis here assumes that the redesigned webpage is not better than the original since $H_0: \pi_1 \geq \pi_2$. Thus we hope to collect data that contradicts this statement so that it can be rejected in favor of the alternative that says the redesign is in fact superior ($H_A: \pi_1 < \pi_2$), and hence worth the expense and effort.

In order to test this hypothesis they randomoized $n_1 = 8,872$ users to the original homepage and $n_2 = 8,642$ users to the redesigned one. In these conditions they observed 280 and 399 conversions, respectively. That is, 280 users in the control condition created accounts while 399 users in the redesign condition created accounts. This sample data is summarized numerically by $\hat{\pi}_1 = 280/8872 = 0.0316$ and $\hat{\pi}_2 = 399/8642 = 0.0462$ which

in practical terms means that 3.16% of users in the control condition created accounts and 4.62% of users in the redesign condition created accounts – corresponding to a 46% increase over the control. We also find $\hat{\pi} = (280 + 399)/(8872 + 8642) = 0.0388$ meaning that the overall account creation rate is 3.88%.

To determine whether the difference in account creation rates between the two conditions is statistically significant, we must formally test the hypothesis by calculating a p-value. To do this, we must first calculate the observed test statistic. Substituting these summaries into equation (2.19) gives

$$t = \frac{0.0316 - 0.0462}{\sqrt{(0.0388)(0.9612) \left(\frac{1}{8872} + \frac{1}{8642} \right)}} = -5.0075.$$

The p-value associated with this test is $P(T \leq -5.0075)$ where $T \sim N(0, 1)$. When calculated this probability is 2.76×10^{-7} , which is effectively 0. In R this probability is calculated using the command `1-pnorm(-5.0075)`.

In order to draw a conclusion, we must compare this value to the significance level $\alpha = 0.05$. Since $2.84 \times 10^{-7} < 0.05$ we reject the null hypothesis in favor of the alternative. In the context of the experiment, this means that the redesigned homepage has a significantly larger likelihood of user account-creation than does the original homepage. Specifically, a 46% increase in account-creation can be expected with the redesigned homepage relative to the original.

2.2.3 Power Analysis and Sample Size Calculations

Here we derive sample size formulae in a manner similar to the development presented in Section 2.1.3 but here we do it in the context of hypothesis tests such as (2.13), (2.14) and (2.15). As in Section 2.1.3 we perform the derivation assuming a two-sided hypothesis is being tested, but we indicate where and how the derivation would change if it were a one-sided hypothesis that was of interest. We also present the derivation in a general manner that does not require equal sample sizes in each condition, and so we assume $n_1 = kn_2$.

As we saw in Section 2.2.1, the null distribution in this scenario is the standard normal distribution – just like it was for the sample size calculations in Section 2.1.3. A convenient consequence of this is that the rejection regions defined in that section are appropriate here as well. The only difference is that we use a different test statistic here. In particular we use the quantity given in equation (2.16). Recall that we preferred not to use this particular equation when actually testing the hypothesis because it required knowing π_1 and π_2 , which are typically not known in practice. However, we note that based on historical data, a data scientist will typically have a good idea of what π_1 is if condition 1 corresponds to the existing product/ platform/ process/ page, etc. Also, when planning the experiment the data scientist will define $\delta = \pi_1 - \pi_2$ to be the effect size (in a manner similar to Section 2.2.1). Thus, with these two pieces of information π_2 can be defined as $\pi_2 = \pi_1 - \delta$, which means we can treat both π_1 and π_2 as known.

As before, we begin by considering the power of the hypothesis test:

$$\begin{aligned}
1 - \beta &= P(\text{Reject } H_0 \mid H_0 \text{ is false}) \\
&= P(T \in R \mid H_0 \text{ is false}) \text{ where } R \text{ is the rejection region} \\
&= P(T \geq z_{\alpha/2} \text{ or } T \leq -z_{\alpha/2} \mid H_0 \text{ is false}) \\
&= P(T \geq z_{\alpha/2} \mid H_0 \text{ is false}) + P(T \leq -z_{\alpha/2} \mid H_0 \text{ is false}) \\
&= P\left(\frac{(\bar{Y}_1 - \bar{Y}_2)}{\sqrt{\frac{\pi_1(1-\pi_1)}{n_1} + \frac{\pi_2(1-\pi_2)}{n_2}}} \geq z_{\alpha/2} \mid H_0 \text{ is false}\right) \\
&\quad + P\left(\frac{(\bar{Y}_1 - \bar{Y}_2)}{\sqrt{\frac{\pi_1(1-\pi_1)}{n_1} + \frac{\pi_2(1-\pi_2)}{n_2}}} \leq -z_{\alpha/2} \mid H_0 \text{ is false}\right)
\end{aligned}$$

If $H_0 : \pi_1 = \pi_2$ were true, and hence $\pi_1 - \pi_2 = 0$ were true, then the ratios in the preceding line would follow a $N(0,1)$ distribution. However, we know that H_0 is false which means that $\pi_1 - \pi_2 = \delta$ for some none-zero δ , and so it is

$$\frac{(\bar{Y}_1 - \bar{Y}_2) - \delta}{\sqrt{\frac{\hat{\pi}_1(1-\hat{\pi}_1)}{n_1} + \frac{\hat{\pi}_2(1-\hat{\pi}_2)}{n_2}}}$$

that follows a $N(0, 1)$ distribution. Let us make this substitution, being sure to replicate what is done on the left side of inequalities on the right. Also note that we no longer need to write “ $|H_0$ is false” since we are now exploiting this fact.

$$\begin{aligned}
1 - \beta &= P \left(\frac{(\bar{Y}_1 - \bar{Y}_2) - \delta}{\sqrt{\frac{\pi_1(1-\pi_1)}{n_1} + \frac{\pi_2(1-\pi_2)}{n_2}}} \geq z_{\alpha/2} - \frac{\delta}{\sqrt{\frac{\pi_1(1-\pi_1)}{n_1} + \frac{\pi_2(1-\pi_2)}{n_2}}} \right) \\
&\quad + P \left(\frac{(\bar{Y}_1 - \bar{Y}_2) - \delta}{\sqrt{\frac{\pi_1(1-\pi_1)}{n_1} + \frac{\pi_2(1-\pi_2)}{n_2}}} \leq -z_{\alpha/2} - \frac{\delta}{\sqrt{\frac{\pi_1(1-\pi_1)}{n_1} + \frac{\pi_2(1-\pi_2)}{n_2}}} \right) \\
&= P \left(Z \geq z_{\alpha/2} - \frac{\delta}{\sqrt{\frac{\pi_1(1-\pi_1)}{n_1} + \frac{\pi_2(1-\pi_2)}{n_2}}} \right) \\
&\quad + P \left(Z \leq -z_{\alpha/2} - \frac{\delta}{\sqrt{\frac{\pi_1(1-\pi_1)}{n_1} + \frac{\pi_2(1-\pi_2)}{n_2}}} \right) \text{ where } Z \sim N(0, 1)
\end{aligned}$$

As in Section 2.1.3 only one of these two terms will dominate, depending on the sign of δ . Assume, without loss of generality, that $\delta > 0$ in which case only the first term remains.

$$1 - \beta = P \left(Z \geq z_{\alpha/2} - \frac{\delta}{\sqrt{\frac{\pi_1(1-\pi_1)}{n_1} + \frac{\pi_2(1-\pi_2)}{n_2}}} \right)$$

Because this probability is equal to $1 - \beta$ we know that $z_{\alpha/2} - \delta / \sqrt{\frac{\pi_1(1-\pi_1)}{n_1} + \frac{\pi_2(1-\pi_2)}{n_2}}$ must be equal to $z_{1-\beta}$, the β^{th} quantile of the standard normal distribution. Thus

$$z_{1-\beta} = z_{\alpha/2} - \delta / \sqrt{\frac{\pi_1(1-\pi_1)}{n_1} + \frac{\pi_2(1-\pi_2)}{n_2}}$$

and we can rearrange this equation solving for the sample size. But first we must substitute $n_1 = kn_2$ so that there is just a single sample size to solve for:

$$z_{1-\beta} = z_{\alpha/2} - \delta / \sqrt{\frac{\pi_1(1-\pi_1)}{kn_2} + \frac{\pi_2(1-\pi_2)}{n_2}} = z_{\alpha/2} - \delta \sqrt{n_2} / \sqrt{\frac{\pi_1(1-\pi_1)}{k} + \pi_2(1-\pi_2)}$$

Solving for n_2 yields:

$$n_2 = \frac{(z_{\alpha/2} - z_{1-\beta})^2 \left[\frac{\pi_1(1-\pi_1)}{k} + \pi_2(1-\pi_2) \right]}{\delta^2} \quad (2.20)$$

and then n_1 is found by computing kn_2 . When equal sample sizes are desired ($k = 1$) each condition receives n units where

$$n = \frac{(z_{\alpha/2} - z_{1-\beta})^2 [\pi_1(1-\pi_1) + \pi_2(1-\pi_2)]}{\delta^2}. \quad (2.21)$$

If it were a one-sided hypothesis being tested, the only difference between the formulae in that setting relative to equations (2.20) and (2.21) above is that in the one-sided case instances of $z_{\alpha/2}$ would be replaced by z_α .

Note that the interactive sample size calculator found at <https://nathaniel-t-stevens.shinyapps.io/SampleSizeCalculator/> can also be used to explore the interdependencies between sample size, significance level, power, and effect size in this setting as well.

2.2.4 Another Way to Think About Comparing Proportions

In order to test the hypotheses (2.13), (2.14) and (2.15) using the Z -test of Section 2.2.1 we required the assumption that the response measurements followed a Bernoulli distribution (i.e., $Y_{ij} \sim \text{BIN}(1, \pi_j)$) and the sample sizes were large enough to ensure the Central Limit Theorem was applicable. Here we describe an equivalent method of testing these hypotheses, but motivated in a slightly different manner. In particular we discuss the **chi-squared test of independence** (also known as Pearson's χ^2 -test).

The chi-squared test of independence is typically used as a test for ‘no association’ between two categorical variables that are summarized in a contingency table. We apply this methodology here to test the independence of the binary outcome (whether a unit performs the action of interest) and the particular condition they are in. If the likelihood of performing the action is the same in each condition (i.e., $\pi_1 = \pi_2$) then the response and conditions

Table 1: 2×2 contingency table for Optimizely’s homepage experiment

		Condition		
		1	2	
Conversion	Yes	280	399	679
	No	8592	8243	16835
		8872	8642	17514

Table 2: A general 2×2 contingency table

		Condition		
		1	2	
Conversion	Yes	$O_{1,1}$	$O_{1,2}$	O_1
	No	$O_{0,1}$	$O_{0,2}$	O_0
		n_1	n_2	$n_1 + n_2$

are not associated. As such, this test is appropriate for evaluating whether $\pi_1 = \pi_2$, $\pi_1 > \pi_2$ or $\pi_1 < \pi_2$. To motivate this, consider the Optimizely data from Section 2.2.2 which have been arranged in a 2×2 contingency table shown in Table 1. When arranged in this fashion we clearly see that there were $n_1 = 8872$ units in condition 1, $n_2 = 8642$ units in condition 2 and there were respectively 280 and 399 conversions in these conditions (and hence 8592 and 8243 non-conversions).

If $\pi_1 = \pi_2 = \pi$ then we would expect the conversion rate in each condition to be the same. An estimate of the pooled conversion rate in this case is $\hat{\pi} = 679/17514 = 0.0388$ since there were 679 conversions in total, and an overall sample size of 17514 users. Thus, we would expect $n_1\hat{\pi} = 8872 \cdot 0.0388 = 344.23$ conversions in condition 1 and $n_2\hat{\pi} = 8642 \cdot 0.0388 = 335.31$ conversions in condition 2. Clearly this is not what we observed, but the chi-squared test formally evaluates if the difference between what was observed and what is expected under the null hypothesis is large enough to be considered *significantly* different.

We formalize this process by considering the general 2×2 contingency table in Table 2, where we let $O_{1,j}$ and $O_{0,j}$ respectively represent the observed number of conversions and non-conversions in condition $j = 1, 2$. Also, O_1 and O_0 represent the overall number of conversions and non-conversions (between both conditions) and so

$$\hat{\pi} = \frac{O_1}{n_1 + n_2} \text{ and } 1 - \hat{\pi} = \frac{O_0}{n_1 + n_2}$$

represent the proportions of units that did or did not convert. As demonstrated above, we use these pooled estimates to calculate the expected number of conversions/non-conversions in each condition. Specifically, we let $E_{1,j}$ and $E_{0,j}$ represent the expected number of conversions and non-conversions in condition $j = 1, 2$ which we calculate as

$$E_{1,j} = n_j \hat{\pi} \text{ and } E_{0,j} = n_j(1 - \hat{\pi}).$$

The χ^2 test statistic compares the observed count in each cell to the corresponding expected count, and is defined as

$$T = \sum_{l=0}^1 \sum_{j=1}^2 \frac{(O_{l,j} - E_{l,j})^2}{E_{l,j}}.$$

Assuming H_0 is true, it can be shown that T approximately follows a χ^2 distribution with $\nu = 1$ degree of freedom (i.e., $T \sim \chi_{(1)}^2$). As a general rule of thumb, this approximation may be very poor unless the observed and expected cell frequencies are all greater than 5.

Then, to draw a conclusion about the hypothesis, we compare the observed value of the test statistic t to the $\chi_{(1)}^2$ distribution. The p-value associated with this test is calculated differently depending on whether H_A is one- or two-sided. These calculations are summarized below:

- $H_A: \pi_1 \neq \pi_2$: p-value = $P(T \geq t)$
- $H_A: \pi_1 > \pi_2$: p-value = $1 - P(T \geq t)/2$
- $H_A: \pi_1 < \pi_2$: p-value = $P(T \geq t)/2$

Although it may not look like it, the value of this χ^2 test statistic will always be the square of value produced by equation (2.19). For this reason, and because squaring a $N(0, 1)$ random variable yields a $\chi_{(1)}^2$ one, the p-values of this χ^2 test will always be identical to those of the Z -test for proportions, and so the conclusions drawn by either method will be the same.

Returning to the Optimizely example, the observed number of conversions in condition 1 and 2 are $O_{1,1} = 280$ and $O_{1,2} = 399$ and the observed number of non-conversions in each condition are $O_{0,1} = 8592$ and $O_{0,2} = 8243$. Using this information we calculate the overall

conversion and non-conversion rates to be $\hat{\pi} = (280 + 399)/(8872 + 8642) = 0.0388$ and $1 - \hat{\pi} = (8592 + 8243)/(8872 + 8642) = 0.9612$. With this we calculate the expected number of conversions in conditions 1 and 2: $E_{1,1} = 343.96$ and $E_{1,2} = 335.04$ and the expected number of non-conversions in conditions 1 and 2: $E_{0,1} = 8528.04$ and $E_{0,2} = 8306.96$. The observed test statistic is then calculated as

$$t = \frac{(280 - 343.96)^2}{343.96} + \frac{(399 - 335.04)^2}{335.04} + \frac{(8592 - 8528.04)^2}{8528.04} + \frac{(8243 - 8306.96)^2}{8306.96} = 25.0755.$$

Then $P(T \geq 25.0755) = 5.52 \times 10^{-7}$, where $T \sim \chi^2_{(1)}$, and the p-values associated with hypotheses (2.13), (2.14) and (2.15) are respectively 5.52×10^{-7} , 0.9999997, and 2.76×10^{-7} . The first p-value suggests that we would reject $H_0: \pi_1 = \pi_2$, suggesting that the conversion rates on the two versions of the homepage are indeed different. The second and third p-values both suggest that $\pi_1 < \pi_2$ is true, indicating that the likelihood of creating an account on the redesigned homepage is higher than on the original version of the homepage.

Note that this test may be implemented in R using the `prop.test()` function.

2.3 The Trouble with Peeking

In Sections 2.1.3 and 2.2.3 we developed sample size calculations to determine the necessary number of units in each condition to ensure the Type I and Type II error rates are held fixed at the predetermined values α and β . However, in practice, the experimentation platform used by a data scientist may provide a dashboard which displays whether, at that current point in time, there is a significant difference between conditions – or that one condition is significantly better than the other.

Often a data scientist may feel external (and/or internal) pressure to stop the experiment when they see this. After all, the results tell us that a winner has been found, right? Wrong. Well, maybe, but by stopping the experiment early you have not observed enough data to be confident in your conclusion. By stopping the experiment you are in effect rejecting the null hypothesis (that the conditions are not different) and so you risk making a Type I error. And by stopping the experiment early the chances you make a Type I error are much higher

than the value α you chose when doing your sample size calculation.

This phenomenon whereby you regularly check the results of the experiment, waiting for a significant result, is known as “peeking”. Peeking is certainly tempting, and depending on your experimentation dashboard, it may be impossible to avoid. In some circumstances, when several metrics are being tracked (in addition to your primary metric of interest) it is in fact a good idea to ‘peek’ to ensure the experiment is not negatively impacting other important metrics.

The problem, however, arises when, as a result of peeking, you decide to end the experiment early. Just because the results suggest a winner or a significant difference at one point in time does not mean that the results won’t change as more data is collected. For instance, I might peek at my experiment now and see that condition 1 is significantly out-performing condition 2. But if I peek again in an hour I might find that condition 2 is significantly out-performing condition 1. Only until you have observed the pre-specified amount of data should you be sure of your conclusions.

To illustrate the dire consequences of peeking and ending an experiment early, consider the following simulated situation. Imagine condition 1 response measurements truly follow a $N(0, 1)$ distribution and condition 2 response measurements also follow a $N(0, 1)$ distribution and an A/B test is performed to decide whether or not $\mu_1 = \mu_2$. In this case the null hypothesis $H_0: \mu_1 = \mu_2$ is true and the data collected should not reject it. However, simply due to chance we may obtain a dataset which leads us to reject H_0 and make a Type I error. However, if the sample sizes n_1 and n_2 were determined so that $\alpha = 0.05$, for example, we would not expect to make this type of error more than 5% of the time.

By repeatedly simulating n_1 and n_2 data points independently from the $N(0, 1)$ distribution, and each time testing the null hypothesis $H_0: \mu_1 = \mu_2$ we can empirically quantify the likelihood of making a Type I error. For illustration we do this $N = 100000$ times, each time with samples of size $n_1 = n_2 = 1000$. In addition to quantifying the Type I error rate if we waited for all $n_1 = n_2 = 1000$ data points to be observed, we also calculate the Type I error rate when the experiment is ended early by peeking at regular intervals.

Here we consider peeking (and ending the experiment if a significant result is indicated) after every successive data point and at intervals of every 2nd, 4th, 5th, 8th, 10th, 20th, 25th, 40th, 50th, 100th, 125th, 200th, 250th, 500th and 1000th data point. This corresponds to peeking 1000 times, 500 times, 250 times, 200 times, 125 times, 100 times, 50 times, 40 times, 25 times, 20 times, 10 times, 8 times, 5 times, 4 times, two times and no peeking at all. For each case in the simulation we peek at the results at the specified interval and end the experiment if the results are statistically significant. We then calculate the Type I error rate as the proportion of the $N = 100000$ times that a Type I error was made. The plot shown in Figure 3 demonstrates how the chances of making a Type I error increase dramatically for increased levels of peeking. Indeed, it is only in the case with no peeking that the Type I error rate is actually equal to 0.05, and with enough peeking, committing a Type I error becomes certain.

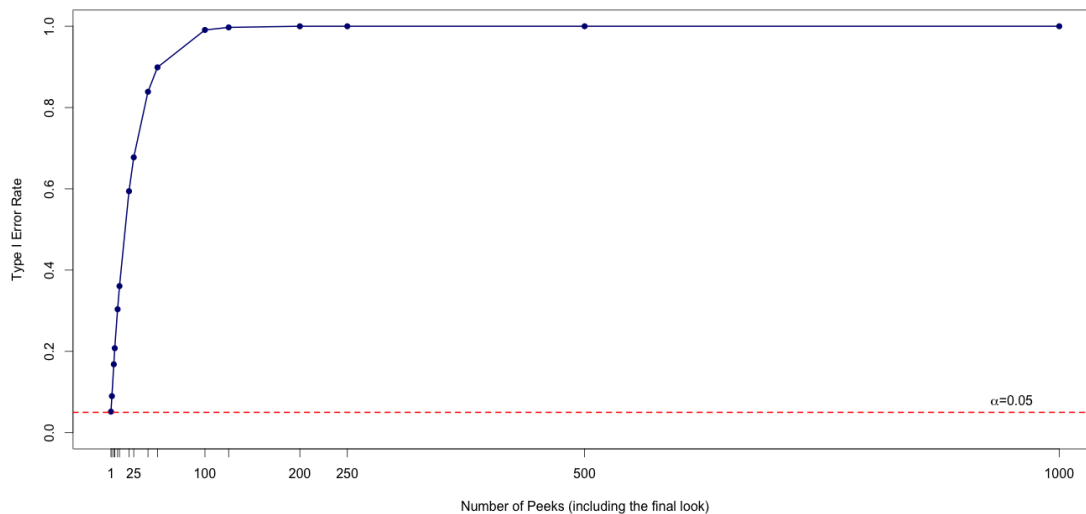


Figure 3: Type I error rate for different levels of peeking.

We note here, and elaborate further later on, that **sequential analysis** and **sequential testing** are important statistical topics/disciplines that are concerned with devising statistically sound methods for performing repeated significance tests as more data becomes available. Essentially, sequential testing corresponds to a host of techniques that allow you to peek and end an experiment early without increasing Type I error rates. However, without adopting one of these techniques, peeking (and ending experiments early) should be avoided

at all costs. We will discuss sequential testing later on in the context of multi-armed bandit experiments.

Experiments With More Than Two Conditions

In the previous chapter we considered the situation in which the experiment contained just two experimental conditions. In the language of designed experiments, this corresponds to the investigation of a single design factor at two levels. We motivated the situation by discussing the canonical A/B test in which two versions of a webpage were compared – one with a red button and the other with a blue button – and we were interested in identifying the ‘winning’ webpage – the one with the button that is clicked most frequently. But what if we want investigate three button colors rather than just two? Or what about 10 button colors?

In many real-life scenarios it is reasonable to believe that a data scientist may be interested in comparing more than just two conditions. For instance, one might be interested in comparing 6 different ads to determine which is most profitable; or, one might be interested in comparing 3 different sign-up promotions to determine which has the highest conversion rate. In general, the question being answered now is: given several options, which is best?

The types of experiments that are used to answer this question are colloquially referred to by data scientists as “A/B/C”, “A/B/C/D”, or more generally, “A/B/n” tests. Formally, these experiments are designed in a very similar manner to A/B tests; a response variable (y) is chosen and some metric of interest θ that summarizes the response measurements is also selected. Recall that this metric may be any statistic that can be calculated from observed data, with various user engagement and conversion metrics being commonly used in practice.

What is different, relative to a traditional A/B test, is the number of levels of the design factor, and hence number of experimental conditions. As before we index experimental conditions by j , but rather than $j = 1, 2$, now we have $j = 1, 2, \dots, m$, where m is the total number of conditions. In this case the metric of interest is calculated in each condition, giving $\hat{\theta}_1, \hat{\theta}_2, \dots, \hat{\theta}_m$, and interest lies in comparing them to determine which condition is optimal. Condition j would be considered optimal if the observed data provided enough evidence to believe $\theta_j > \theta_k$ for all $k \neq j$ (when maximizing the metric is important) or $\theta_j < \theta_k$ for all $k \neq j$ (when minimizing the metric is important).

In this chapter we describe the statistical tests that are used to draw conclusions of this sort, and we discuss practical and statistical problems that must be considered in this situation. Like the previous chapter we consider the comparison of means and the comparison of proportions.

3.1 Comparing Means in Multiple Conditions

As in Section 2.1, we assume that our response variable follows a normal distribution and we assume that the mean of the distribution depends on the condition in which the measurements were taken, and that the variance is the same across all conditions. Mathematically, we assume $Y_{ij} \sim N(\mu_j, \sigma^2)$ for $i = 1, 2, \dots, n_j$ and $j = 1, 2, \dots, m$. Thus, the only difference between these distributions (if there is a difference) is in their means. As such, formal hypothesis tests in this scenario concern only the μ_j 's. While interest ultimately lies in finding the condition with the highest (or smallest) μ_j , a common (and sensible) starting point is to decide whether there is a difference at all between the conditions. To answer this question formally, the following hypothesis is tested.

$$H_0: \mu_1 = \mu_2 = \dots = \mu_m \text{ vs. } H_A: \mu_j \neq \mu_k \text{ for some } k \neq j \quad (3.1)$$

Failing to reject H_0 means that the expected response does not differ significantly from one condition to another, and so no single condition is optimal. However, if the observed data provide enough evidence to reject H_0 , then we would conclude that the expected response

in at least one of the conditions is not the same as the others. Follow-up hypothesis tests can then be used to determine which condition(s) is (are) optimal. These follow-up tests are typically performed in a pairwise manner, comparing a given condition to each of the other conditions. The two-sample methods discussed in Section 2.1 are useful for this task. However, it is important to note that when doing multiple comparisons and hence testing a series of hypothesis tests, the overall Type I error rate becomes inflated, and so modifications to the testing procedure must be made. We discuss this further in Section 3.3.

In the next subsection we discuss how to formally test hypothesis (3.1). As we will see, this test can be performed using the **F-test for overall significance** in an appropriately defined linear regression model. For a primer on linear regression, see Appendix A.3.

3.1.1 The F -test for Overall Significance in a Linear Regression

The “appropriately defined linear regression model” in this situation is one in which the response variable depends on $m - 1$ indicator variables where, for example, the indicator variables may be defined as

$$x_{ij} = \begin{cases} 1, & \text{if unit } i \text{ is in condition } j \\ 0, & \text{otherwise} \end{cases}$$

for $j = 1, 2, \dots, m - 1$. For a particular unit i , we adopt the model

$$Y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_{m-1} x_{i,m-1} + \epsilon_i$$

where Y_i is the response observation for unit $i = 1, 2, \dots, N = \sum_{j=1}^m n_j$ and ϵ_i is the corresponding random error term assumed to follow a $N(0, \sigma^2)$ distribution.

In this model the β 's are unknown parameters which we interpret in the following manner. The intercept β_0 is the expected response when each of the indicator variables is equal to zero: $E[Y_i | x_{i1} = x_{i2} = \dots = x_{i,m-1} = 0] = \beta_0$. By design, if all of the indicator variables are equal to zero, this means that unit i was in condition m . Thus β_0 is the expected response

in condition m .

Similarly, since $E[Y_i|x_{ij} = 1] = \beta_0 + \beta_j$, we define $\beta_0 + \beta_j$ to be the expected response in condition $j = 1, 2, \dots, m - 1$, and interpret β_j as being the expected change in response in condition $j = 1, 2, \dots, m - 1$ relative to in condition m . It is practically useful to treat condition m as the ‘control’ condition, if there is one, since it represents the baseline against which all other conditions are compared.

Thus,

$$\begin{aligned}\mu_1 &= \beta_0 + \beta_1 \\ \mu_2 &= \beta_0 + \beta_2 \\ &\vdots \\ \mu_{m-1} &= \beta_0 + \beta_{m-1} \\ \mu_m &= \beta_0\end{aligned}$$

As can be seen, H_0 in (3.1) is true if and only if $\beta_1 = \beta_2 = \dots = \beta_m = 0$. Thus testing (3.1) is equivalent to testing

$$H_0: \beta_1 = \beta_2 = \dots = \beta_m = 0 \text{ vs. } H_A: \beta_j \neq 0 \text{ for some } j$$

in the context of the linear regression model above. Such a test is known as the F -test for overall significance in a linear regression. The test statistic is defined to be the ratio of the regression mean squares (MSR) to the mean squared error (MSE) that are associated with a standard regression-based analysis of variance (ANOVA):

$$t = \frac{MSR}{MSE}.$$

Note that MSE is an estimate of σ^2 , as described in Appendix A.3, and MSR is related to the MSE of the *reduced model* that assumes H_0 is true (i.e., $\beta_1 = \beta_2 = \dots = \beta_m = 0$).

Assuming the null hypothesis is true, this test statistic should look as if it comes from an F -distribution with $\nu_1 = m - 1$ and $\nu_2 = N - m$ degrees of freedom. The p-value associated

with this test is calculated as $p\text{-value} = P(T \geq t)$ where $T \sim F(m - 1, N - m)$ and is commonly displayed in regression summaries provided by statistical software. For example, the `summary()` of an `lm()` object in R provides the results of this test. We illustrate the use of this test with an example in the next subsection.

3.1.2 Example: Candy Crush Boosters

Candy Crush is experimenting with three different versions of in-game “boosters”: the lollipop hammer, the jelly fish, and the color bomb. Users are randomized to one of these three conditions ($n_1 = 121$, $n_2 = 135$, $n_3 = 117$) and they receive (for free) 5 boosters corresponding to their condition. Interest lies in evaluating the effect of these different boosters on the length of time a user plays the game. Let μ_j represent the average length of game play (in minutes) associated with booster condition $j = 1, 2, 3$. While interest lies in finding the condition associated with the longest average length of game play, here we first rule out the possibility that booster type does not influence the length of game play (i.e., $\mu_1 = \mu_2 = \mu_3$). In order to do this we fit the linear regression model

$$Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \epsilon$$

where the x ’s are indicator variables indicating whether a particular value of the response was observed in the jelly fish or color bomb conditions. By using the `lm()` function in R we obtain the following output

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	5.01281	0.08664	57.859	<2e-16 ***
factor(booster)2	1.17528	0.11931	9.851	<2e-16 ***
factor(booster)3	4.88279	0.12357	39.515	<2e-16 ***

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

Residual standard error: 0.953 on 370 degrees of freedom

Multiple R-squared: 0.8216, Adjusted R-squared: 0.8206
F-statistic: 851.9 on 2 and 370 DF, p-value: < 2.2e-16

From this output we see that $\hat{\beta}_0 = 5.0128$, $\hat{\beta}_1 = 1.1753$ and $\hat{\beta}_2 = 4.8828$ indicating the average length of game play is estimated to be $\hat{\mu}_1 = \hat{\beta}_0 = 5.0128$ minutes in the lollipop hammer condition, $\hat{\mu}_2 = \hat{\beta}_0 + \hat{\beta}_1 = 5.0128 + 1.1753 = 6.1881$ minutes in the jelly fish condition, and $\hat{\mu}_3 = \hat{\beta}_0 + \hat{\beta}_2 = 5.0128 + 4.8828 = 9.8956$ minutes in the color bomb condition. These estimates suggest that the average length of game play differs depending on which booster condition a user is in. To formally draw this conclusion we perform the F -test of overall significance of the regression. The output shown above indicates that the observed test statistic is calculated to be $t = 851.9$ and the p-value $= P(T \geq 851.9)$ is less than 2.2×10^{-16} , where T follows an F -distribution with $\nu_1 = 2$ and $\nu_2 = 370$ degrees of freedom. Such a small p-value provides very strong evidence against H_0 and so we conclude that the average length of game play is not the same for each of the boosters. To determine which booster is optimal – the one that maximizes game play duration – we must use a series of pairwise t -tests. This is left as an exercise for the reader.

3.2 Comparing Proportions in Multiple Conditions

Like in Section 2.2, we assume that our response variable is binary:

$$Y_{ij} = \begin{cases} 1, & \text{if unit } i \text{ in condition } j \text{ performs an action of interest} \\ 0, & \text{if unit } i \text{ in condition } j \text{ does not perform an action of interest} \end{cases}$$

for $i = 1, 2, \dots, n_j$ and $j = 1, 2, \dots, m$. As before we define $\pi_j = P(Y_{ij} = 1)$ to be the probability that a unit in condition j performs the action – whether this means that they click a button, open an email, create an account, etc. Of interest, then, is a comparison of the likelihood that the action is performed across all conditions, with the ultimate goal of finding the condition with the highest (or smallest – whichever corresponds to optimal) π_j . While this is the ultimate goal, a sensible first step is to decide whether there is a difference between the conditions at all. In order to formally make this decision, the following hypothesis is

tested.

$$H_0: \pi_1 = \pi_2 = \dots = \pi_m \text{ vs. } H_A: \pi_j \neq \pi_k \text{ for some } k \neq j \quad (3.2)$$

Failing to reject H_0 means that the action of interest is no more probable in one condition than any other, and so no single condition is optimal. However, if the observed data provide enough evidence to reject H_0 , then we would conclude that there is at least one condition in which units behave differently. Follow-up hypothesis tests can then be used to determine which condition(s) is (are) optimal. These follow-up tests are typically performed in a pairwise manner, comparing a given condition to each of the other conditions. The two-sample methods discussed in Section 2.2 are useful for this task. However, we remark again that performing multiple comparisons can lead to an increased Type I error rate, which we discuss further in Section 3.3.

In the next subsection we discuss how to formally test hypothesis (3.2). As we will see, the χ^2 test from Section 2.2.4 generalizes to the comparison of any number of conditions and so we will apply it again in this scenario.

3.2.1 The Chi-squared Test of Independence

In Section 2.2.4 we introduced the χ^2 test of independence as a means to evaluate whether $\pi_1 = \pi_2$, $\pi_1 > \pi_2$, or $\pi_1 < \pi_2$ in the context of two experimental conditions. However, we noted that the test was more generally used as a test of ‘no association’ between two categorical variables, and so we could think of it as a test of ‘no association’ between the binary outcome (whether a unit performs the action of interest) and the particular condition they are in. However, before we considered just two conditions – but the test itself imposes no such restriction; here we consider the test more generally, in the context of m experimental conditions.

As before we are interested in comparing observed and expected frequencies of each outcome in each condition. The information associated with this test can be summarized in a $2 \times m$ contingency table, where rows correspond to the binary outcome, conversion, and the columns correspond to the different conditions. The value in the $(l, j)^{th}$ cell of this table,

Table 3: A general $2 \times m$ contingency table
Condition

Conversion		1	2	\dots	m	
	Yes	$O_{1,1}$	$O_{1,2}$	\dots	$O_{1,m}$	O_1
	No	$O_{0,1}$	$O_{0,2}$	\dots	$O_{0,m}$	O_0
		n_1	n_2	\dots	n_m	$N = \sum_{j=1}^m n_j$

denoted $O_{l,j}$, corresponds to the observed number of conversions ($l = 1$) or non-conversions ($l = 0$) in condition $j = 1, 2, \dots, m$. An example of such a table is shown in Table 3.

As in the 2×2 case, each of these observed frequencies is contrasted with an expected frequency where $E_{1,j}$ is the expected number of conversions in condition j and $E_{0,j}$ is the expected number of non-conversion in condition j . These expected frequencies are found by multiplying condition j 's sample size by the pooled conversion and non-conversion rates:

$$E_{1,j} = n_j \hat{\pi} \text{ and } E_{0,j} = n_j(1 - \hat{\pi})$$

where

$$\hat{\pi} = \frac{O_1}{N} \text{ and } (1 - \hat{\pi}) = \frac{O_0}{N}$$

are the sample estimates of homogenous probabilities of conversion and non-conversion.

The test statistic for this test is defined exactly as it was in the 2×2 case except that now we are summing over more cells and the null distribution has a different number of degrees of freedom. In particular, the test statistic is

$$T = \sum_{l=0}^1 \sum_{j=1}^m \frac{(O_{l,j} - E_{l,j})^2}{E_{l,j}}$$

where, if H_0 is true, an observed value, t , should look as if it comes from a χ^2 distribution with $\nu = m - 1$ degrees of freedom. The p-value associated with this test is calculated as $\text{p-value} = P(T \geq t)$ where $T \sim \chi_{(m-1)}^2$. As in the 2×2 case, this test can be carried out automatically using the `prop.test()` function in R. We illustrate the use of this test with an example in the next subsection.

Table 4: A 2×5 observed contingency table for the Nike example

		Condition				
		1	2	3	4	5
View	Yes	160	95	141	293	197
	No	4854	4876	4889	4714	4783
		5014	4971	5030	5007	4980

Table 5: A 2×5 expected contingency table for the Nike example

		Condition				
		1	2	3	4	5
View	Yes	177.68	176.16	178.25	177.43	176.48
	No	4836.32	4794.84	4851.75	4829.57	4803.52
		5014	4971	5030	5007	4980

3.2.2 Example: Nike SB Video Ads

Suppose that Nike is running an ad campaign for Nike SB, their skateboarding division, and the campaign involves $m = 5$ different video ads the are being shown in Facebook newsfeeds. A video ad is ‘viewed’ if it is watched for longer than 3 seconds, and interest lies in determining which ad is most popular and hence most profitable by comparing the viewing rates of the five different videos. Each of these 5 videos is shown to $n_1 = 5014$, $n_2 = 4971$, $n_3 = 5030$, $n_4 = 5007$, and $n_5 = 4980$ users and in each condition the videos are viewed 160, 95, 141, 293 and 197 times, respectively, yielding watch rates of $\hat{\pi}_1 = 0.0319$, $\hat{\pi}_2 = 0.0191$, $\hat{\pi}_3 = 0.0280$, $\hat{\pi}_4 = 0.0585$, and $\hat{\pi}_5 = 0.0396$.

Based on these estimates it would suggest that not all of the videos are equally popular, but to formally decide this we will conduct a χ^2 -test. The 2×5 contingency table for these data are shown in Table 4.

The expected cell frequencies are found by multiplying n_j by $\hat{\pi}$ and $(1 - \hat{\pi})$, $j = 1, 2, 3, 4, 5$, where $\hat{\pi}$ is calculated using these data to be $\hat{\pi} = 886/25002 = 0.0354$. Table 5 displays these frequencies. The observed test statistic for these data is calculated to be

$$t = \sum_{l=0}^1 \sum_{j=1}^5 \frac{(O_{l,j} - E_{l,j})^2}{E_{l,j}} = 129.1761.$$

The p-value, then, is $P(T \geq 129.1761) = 5.84 \times 10^{-27}$ where $T \sim \chi^2_{(4)}$. Such a small p-value provides very strong evidence against H_0 in this case, and so we conclude that the likelihood that someone ‘views’ a video is not the same for all of the videos. To determine which video is optimal – the one with the highest likelihood of viewing – we must use a series of pairwise Z -tests or χ^2 -tests. This is left as an exercise for the reader.

3.3 The Problem of Multiple Comparisons

As the examples in Sections 3.1.2 and 3.2.2 illustrate, the hypothesis of overall equality (see e.g., (3.1) or (3.2)) is often rejected. In these situations, a series of follow-up pairwise comparisons are necessary to determine which condition(s) is (are) optimal. From a practical standpoint, we are already armed with the statistical machinery to do this; we need only perform several two-sample t -tests, Z -tests, χ^2 -squared tests or randomization tests (whatever the situation calls for). However, when doing multiple comparisons like this, it is important to recognize that if each individual test has a Type I error rate of α , the overall Type I error rate associated with this family of tests, is much larger than α .

This problem – where a series of independent hypothesis tests lead to an inflated family-wise error rate – is known as the **multiple comparison** or **multiple testing problem**. It can be shown that if a family of k hypothesis tests are performed, each with significance level α , the family-wise error rate (the probability of making a Type I error in any of these

tests) is $1 - (1 - \alpha)^k$. To see this:

$$\begin{aligned}
P(\text{Type I Error}) &= 1 - P(\text{No Type I Error}) \\
&= 1 - P(\{\text{No Type I Error on test 1}\} \text{ and } \{\text{No Type I Error on test 2}\} \text{ and } \\
&\quad \dots \text{ and } \{\text{No Type I Error on test } k\}) \\
&= 1 - P(\{\text{No Type I Error on test 1}\} \cap \{\text{No Type I Error on test 2}\} \cap \\
&\quad \dots \cap \{\text{No Type I Error on test } k\}) \\
&= 1 - \prod_{i=1}^k P(\text{No Type I Error on test } i) \\
&= 1 - \prod_{i=1}^k (1 - \alpha) \\
&= 1 - (1 - \alpha)^k
\end{aligned}$$

Figure 4 illustrates the dependence of this family-wise error rate on the number of pairwise comparisons, k . As we can see, as k increases so also does the error rate. In the limit (i.e., as $k \rightarrow \infty$) this error rate goes to 1, and so as the the number of pairwise comparisons increases it becomes certain that a Type I error will have been made somewhere. In practice, a common value of k is $\binom{m}{2}$: the number of pairwise comparisons necessary to compare each condition to every other condition. Supposing the experiment consists of $m = 5$ experimental conditions, then $k = \binom{5}{2} = 10$ which, if the significance level on each test is $\alpha = 0.05$, results in a family-wise error rate of 0.4013 – much higher than the Type I error rate we are comfortable with.

In order to combat this, a variety of statistical approaches have been developed. Here we consider the simplest and most commonly used: the **Bonferroni correction** (Dunnett, 1955). In order to keep the family-wise error rate maintained at α , the Bonferroni correction involves performing each of the k hypothesis tests at a significance level of α/k . Doing so yields a family-wise error rate of $1 - (1 - \frac{\alpha}{k})^k$ which, for typical values of α in the range $(0, 0.05]$ is approximately equal to α . Figure 5 illustrates this; across a range of values for k the family-wise error is held fixed at the specified level of significance α .




Figure 4: Family-wise error rate versus the number of pairwise comparisons, k .




Figure 5: Illustration of the Bonferroni correction for selected values of α .

Furthermore, we can see that as k gets asymptotically large the family-wise error rate no longer approaches 1. Instead

$$\lim_{k \rightarrow \infty} 1 - \left(1 - \frac{\alpha}{k}\right)^k = e^{-\alpha},$$

which for typical values of α in the range $(0, 0.05]$ is approximately equal to α . Figure 6 illustrates this.




Figure 6: Illustration of the Bonferroni correction for asymptotically large k .

So what does this mean for sample size calculations and power analyses? In Sections 2.1.3 and 2.2.3 we showed that sample size formulae could be derived which accounted for the desired power and significance level of the test. However, this did not account for the multiple comparison problem. If, at the time of designing the experiment, you know that you intend to do k pairwise hypothesis tests in order to find a ‘winning’ condition, then the significance level you use in your sample size calculations should be adjusted to account for this. In particular, if applying the Bonferroni correction, one should use α/k as the significance level in these calculations, if a family-wise error rate of α is to be maintained.

Multi-armed Bandit Experiments

The broad goal of experimentation in the world of data science is optimization; given two or more conditions, which one(s) is (are) best? Throughout these notes we have considered the design and analysis of many such experiments. In all cases the experimental conditions have been defined by the unique combinations of one or more factors. The manner in which these experiments are analyzed depends upon the type of metric (θ) being compared across conditions and by consequence the data type of the response variable. But by whatever method, the goal has always been the same: find the optimal condition. Mathematically: given observed data from each of m conditions, interest lies in finding the optimal θ_j among $\theta_1, \theta_2, \dots, \theta_m$.

The comparison of $m \geq 2$ conditions, where the goal is to find the optimal condition, may be thought of as a **multi-armed bandit** problem. Slot machines are often referred to colloquially as “one-armed bandits” because you pull their arm and they take your money. A row of slot machines then is colloquially referred to as a “multi-armed bandit” since there are several arms that, when pulled, will take your money. The multi-armed bandit problem concerns finding the slot machine (i.e., the arm) with the highest expected reward. One can imagine sitting at a row of slot machines in practice and repeatedly playing each machine trying to empirically find the optimal machine. Many mathematical and algorithmic solutions have been proposed, though their solutions are only approximate. Finding an exact solution (i.e, an algorithm that will find the optimal arm with certainty) is very difficult ².

The general goal of finding an optimal condition (i.e., arm) is exactly what we have been doing! However, the approach we have been taking is what some might call the “classical”

²Whittle (1979) facetiously remarked that when bandit problems were first formulated during World War II, allied forces suggested that the problem be “dropped over Germany as the ultimate instrument of intellectual sabotage”

approach; in our experiments we collect a certain amount of data in accordance with Type I and Type II error constraints, and once the data have been observed, we conduct a hypothesis test. Typical multi-armed bandit solutions differ from classical experiments largely with respect to the **exploration-exploitation** trade off. In this chapter we will define this trade off and discuss the general multi-armed bandit solution as well as some specific solutions. References for further reading will also be provided.

4.1 The General Multi-armed Bandit Problem

When searching for an optimal condition we must **explore** all of the options, and once the optimal one is found, we would like to **exploit** it. In the context of slot machines, you'd like to play all of the machines in order to find the optimal one. Then, once you've found the optimal one you simply play it repeatedly, thus maximizing your rewards. In the context of experiments, you want to observe experimental units in every condition in order to find the optimal one. Then, once you've found it you simply direct all units to that condition, thus optimizing your metric of interest θ .

In the classical approach, the experiment corresponds to a phase of 100% exploration, which is followed by 100% exploitation. Figure 7 depicts this exploration-exploitation trade off in the case of classical experiments. Here the experiment provides the opportunity for only exploration. The optimal condition is not exploited until the experiment has concluded.




Figure 7: Exploration-Exploitation trade off in classical experiments.

Multi-armed bandit experiments, on the other hand, are typified by periods of both exploration and exploitation during the experiment itself. Figure 8 depicts the exploration-exploitation trade off in these types of experiments. Whereas experimental units are allocated to conditions at a fixed rate in the classical experiment, multi-armed bandit experiments do not do this. Instead, at regular intervals the proportion of units allocated to each condition is updated to reflect the performance of each condition observed to that point in time.




Figure 8: Exploration-Exploitation trade off in multi-armed bandit experiments.

Units are assigned to experimental conditions with equal frequency initially, but in multi-armed bandit experiments high performing conditions receive more units than lower performing conditions. As such, the multi-armed bandit experiment both explores each of the conditions, but also exploits high performing conditions, even before the experiment has concluded. The proposed advantage of such an approach is that optimal conditions are supposedly found more quickly than with the classical experimental approach. A second advantage is multi-armed bandit experiments assign fewer units to under performing conditions. Consequently this adaptive assignment of units to conditions reduces the opportunity cost associated with experimenting with a risky or inferior condition; rather than assigning many units to such a condition (in accordance with a power analysis), extra units are assigned to the high performing conditions.

However, multi-armed bandit experiments are not always appropriate. Such experiments completely ignore the the impact of Type I and Type II error; sample sizes are not considered here and experiments are ended far sooner that within the classical paradigm. By adopting this methodology, one must believe that the consequences associated with these types of er-

ror are of no practical importance. Thus, if avoiding Type I and Type II error are important, then the classical experimental approach is more appropriate. A second situation for which multi-armed bandit experiments are not appropriate is when there is a delay between assigning a unit to an experimental condition and observing their response value. The adaptive assignment of units to conditions requires that response observations are observed almost instantaneously – in order to determine the assignment strategy in the next time period, we require an accurate estimate of the performance of each condition in the current time period; delayed response observations jeopardize this. A common place where such a delay can be expected is email advertising: experimental units may receive different versions of an email promotion, but there may be a significant delay between the time of email receipt and the time of conversion. This type of experiment is not well-suited for the adaptive assignment strategy that accompanies most multi-armed bandit experiments.

In the next section we discuss four specific adaptive assignment strategies, and emphasize one in particular that works well in practice (when an adaptive assignment strategy is appropriate).

4.2 Some Specific Multi-armed Bandit Solutions

Here we describe four different strategies for assigning experimental units to conditions that vary with respect to the exploration-exploitation spectrum.

4.2.1 Equal Allocation

As the name suggests, each condition is allocated experimental units in equal proportions for the duration of the experiment. This approach is not adaptive and in fact corresponds to the classical experimental paradigm. As we have seen, this represents 100% exploration during the experiment, with no exploitation. Detractors would say that this is inefficient and that the optimal condition can be found more quickly with an adaptive allocation strategy.

4.2.2 Greedy Approach

After an initial period of equally allocating experimental units to conditions, the allocation of units in every subsequent time period is based on the highest performing condition. In other words, at a given time point, every experimental unit is assigned to the ‘best’ condition as determined by the data observed up to that point in time. This approach is sometimes referred to as “play-the-winner” and it represents 100% exploitation. Due to the greedy nature of this algorithm, this approach may do a poor job at maximizing rewards as it does not adequately explore other conditions.

4.2.3 Epsilon-Greedy Approach

The epsilon-greedy approach is a hybrid of the previous two that forces both exploration and exploitation. In this approach unit allocation is performed via the greedy approach with probability $1 - \epsilon$ and equal allocation with probability ϵ . Thus for any experimental unit, one can think of a binary random number being generated such that a 1 indicates greedy allocation, in which case the unit is assigned to the current ‘best’ condition. However, if the random number is a 0 then equal allocation is implemented and the unit is assigned at random to any one of the m conditions. Notice that the equal allocation ($\epsilon = 1$) and greedy allocation ($\epsilon = 0$) approaches of Sections 4.2.1 and 4.2.2 are special cases of the epsilon-greedy approach. Choices of $0 < \epsilon < 1$ determine the desired balance of exploration and exploitation and is determined by the experimenter. Note that a potential drawback of the epsilon-greedy approach is that it will continue to explore even once an optimal condition has been found. However, *epsilon-decreasing* methods have been developed to combat this problem.

4.2.4 Randomized Probability Matching

Randomized probability matching (RPM) is a **Bayesian** approach to the adaptive allocation problem that relies on Thompson sampling (Scott, 2010). With this approach, the

probability that θ_j is optimal is calculated at every time point for each θ_j , $j = 1, 2, \dots, m$. This probability calculation is based on the joint posterior distribution of $(\theta_1, \theta_2, \dots, \theta_m)$ given the data observed up to that point in time. These m optimality probabilities (one for each condition) are used as allocation weights in the next round of allocation.

As an example, consider a standard A/B (two-condition) test where at the current point in time condition A has a 0.73 probability of being the superior condition, and condition B has a 0.27 probability of being the superior condition. Then, in the next round of assigning units to conditions, condition A would receive 73% of the experimental units and condition B would receive 27% of them. The experiment continues in this way until the probability of optimality dominates for one of the conditions, while the other probabilities of superiority tend to zero.

Figure 9 depicts the evolution of optimality probabilities in a simulated experiment with $m = 2$ conditions where condition 2 is truly superior with a conversion rate of 0.05, and condition 1 is truly inferior with a conversion rate of 0.04. In this figure we see that at the beginning of the experiment, when units are allocated equally to the two conditions, each is equally likely to be the optimal condition. However, as we observe data these probabilities are updated; early on (when sample sizes are still relatively small) there is a lot of uncertainty regarding which condition is the optimal one. However, as time (and data) accumulate it becomes clear that condition 2 is optimal – we see that its probability of optimality tends to 1 while the probability that condition 1 is optimal tends to 0.

Figure 10 depicts the evolution of optimality probabilities in a second simulated experiment – this time with $m = 6$ conditions. In this case condition 2 is optimal (with a conversion rate of 0.05) and condition 3 is second best (with a conversion rate of 0.045). The other four conditions have conversion rates of 0.04 or less. Again we see that early on there is a lot of uncertainty about which condition is best, but as more data is collected it becomes more and more clear that condition 2 is best. Note that the algorithm learned early that conditions 1, 4, 5 and 6 were suboptimal, but it took slightly longer to distinguish between conditions 2 and 3, whose conversion rates are very similar.

Note that if the optimality probability does not dominate for one condition in particular,




Figure 9: Simulated evolution of optimality probabilities in an experiment with $m = 2$ conditions. Source: <https://support.google.com/analytics/answer/2844870?hl=en>




Figure 10: Simulated evolution of optimality probabilities in an experiment with $m = 6$ conditions. Source: <https://support.google.com/analytics/answer/2844870?hl=en>

it suggests that multiple conditions are equally optimal. If, among m conditions, h of them were equally optimal, this result would be manifested as the h optimality probabilities stabilizing at roughly $1/h$.

[Scott \(2010\)](#) has shown that RPM does a better job at balancing exploration and exploitation than other allocation strategies, and is able to more quickly find the optimal condition. Although this approach works well in practice, its implementation is less straightforward than other strategies (like the epsilon-greedy strategy, for example) because it relies on an understanding of Bayesian statistics and the ability to sample from a posterior distribution obtained from Markov chain Monte Carlo (MCMC) simulation. That said, for automated implementation, the RPM solution to multi-armed bandit experiments is the default allocation strategy in Google Analytics' experimentation platform.

For a thorough overview of the multi-armed bandit problem and a comparison of the performance of several algorithmic solutions, see [Kuleshov and Precup \(2000\)](#).

Experiments With Multiple Factors

Thus far we have considered experiments with two or more experimental conditions. But no matter the number of conditions, they have always been derived from the levels of a single design factor. However, in most practical circumstances, there might be several factors that are expected to impact the response variable. As such, in this chapter, we consider experimentation with multiple design factors – sometimes referred to as ‘multivariate experiments’. In previous chapters we have motivated the discussion by considering a toy example in which an experiment is performed to determine the influence of a button’s color on the likelihood that the button gets clicked. But what if you also wish to investigate the influence of the button’s size or location, or the button’s message, on the likelihood that it gets clicked? Here we describe how to design and analyze experiments that efficiently investigate multiple factors.

Of particular interest is finding the combination of factor levels that optimize the response variable of interest. Returning to the button example, is it a large red button that says “Go!” that is most likely to get clicked, or is it a medium-sized green one that says “Submit”? Similar to Chapter 3, experiments in this setting involve multiple conditions where the goal is to find the optimal one(s), but the difference here lies in how the conditions arise. In the next section we describe the **factorial** approach and argue that it is superior to simpler but less informative alternatives, and later in the chapter we demonstrate how to design and statistically analyze factorial experiments.

5.1 The Factorial Approach

The key to multi-factor experiments is to efficiently investigate different combinations of the factor levels so as to identify an optimal combination. A common (and simple) method of doing this is the **one-factor-at-a-time** approach in which a sequence of experiments is performed, each with just one factor being varied. Such an approach is manifested as a sequence of single-factor multi-level experiments in which the winning level in a given experiment is retained in future experiments while some other factor is manipulated.

As an example, consider the change Twitter implemented a few years ago in which “favorites” (expressed as stars) were replaced by “likes” (expressed as hearts) and the color of the heart was chosen to be red (as opposed to say, yellow). These decisions likely came about through a series of experiments in which the icon’s shape and color were changed. The factors, then, are shape and color with levels respectively given by $\{star, heart\}$ and $\{red, yellow\}$. This represents the simplest possible multi-factor situation: two factors, each with two levels. To investigate this using the one-factor-at-a-time approach an initial ‘shape’ experiment would be conducted (i.e., an A/B test to determine which shape was most popular). Then, supposing the heart was the winning shape, a follow-up ‘color’ experiment to determine the optimal color would be performed (i.e., a second A/B test to determine whether red hearts or yellow hearts are preferable).

While sequential and iterative learning are important tenets of experimentation, a single, more efficient and more informative, experiment could have been performed in place of this less efficient sequence of two experiments. Note that within the one-factor-at-a-time approach, a red star was never investigated. What if the red star truly out-performs the three shape-color combinations that were investigated? The only way to determine this would have been to formally run a *red star* condition.

The factorial approach to multi-factor experimentation is protected against this short-coming – potentially optimal conditions are not missed because *every* combination of factor levels is considered. A factorial approach to the Twitter experiment would consist of four conditions stemming from the four possible shape-color combinations: yellow star, red star,

yellow heart, red heart. And assuming the one-factor-at-a-time approach consisted of two A/B tests, a total of four conditions were investigated anyway. So the one-factor-at-a-time approach takes as many conditions as the factorial approach! This represents a missed opportunity to simply investigate all possible combinations for which there is no loss of information.

Note that as the number of factors and number of levels increase beyond two, the number of conditions associated with a factorial experiment will necessarily be larger than with the one-factor-at-a-time approach. However, one does not risk not finding an optimal combination with this approach; a factorial experiment is the most efficient method of fully investigating multiple factors as it allows for the quantification of both **main effects** and **interaction effects**. The main effect of factor A (let's call it), represents the change in the response variable produced by a change in that factor. If the factor has a large influence on the response variable, then the magnitude of its main effect will be large. However, sometimes the main effect of one factor, A, depends on the level of another factor, B. In this case we say that factors A and B **interact** and a large interaction effect will lead to the main effect of A being very different for different levels of B. We will make these definitions and ideas more precise in Section 5.3. But first, in Section 5.2, we discuss the practical considerations that need to be made when designing a factorial experiment.

5.2 Designing a Factorial Experiment

Conceptually, the design of a factorial experiment is simple: select a number of design factors you believe influence the response, select the levels you'd like to experiment with for each of these design factors, and then define the experimental conditions to be each of the possible combinations of these factors' levels. For instance, suppose you are interested in investigating the factors labeled '1', '2', '3', and '4' which have m_1 , m_2 , m_3 and m_4 levels, respectively. The full factorial experiment that investigates these factors consists of $m_1m_2m_3m_4$ experimental conditions, representing all possible combinations of the four factors' levels. In general, a full factorial experiment with k factors requires $M = m_1m_2 \cdots m_k$ conditions.

Clearly, as the number of factors and/or levels increases, the number of experimental conditions gets large. This is the primary drawback of factorial experiments: they get big, quickly. So as not to design an experiment that is unmanageably large, careful thought should be given to both the selection of design factors and their associated levels. In particular, it would be a waste of effort to investigate factors that are highly correlated, or factors that are difficult to manipulate outside the confines of a small controlled experiment. It would also be a waste of effort to investigate several levels of a factor that is uninfluential. In general, when choosing factors and levels in the context of a factorial experiment, these choices should be made with the “keep it simple” principle in mind. A special class of factorial experiments is the one in which each factor is investigated at just *two* levels. Such a factorial experiment minimizes the number of experimental conditions required to investigate multiple factors, and as such is useful for **screening** factors to determine which ones are influential and which ones are not. Further experimentation can then be performed with the influential factors at more than two levels. We discuss these so-called “ 2^k factorial experiments” in more detail in the next chapter.

Once the factors, levels, and hence experimental conditions have been established, experimental units must be randomized to each of the M conditions. The number of experimental units to assign to each condition, denoted n_j ($j = 1, 2, \dots, M$), may differ from one condition to the next, but we may also require that they be the same (i.e., $n_1 = n_2 = \dots = n_M = n$). The number of units assigned to each condition (whether the same or different in each condition) can be determined using the pairwise sample size calculations discussed in Chapter 2. In general this type of experiment can have two goals:

- (1) Identify which combination of factor levels (i.e., which condition) is optimal
- (2) Identify which factors are influential

In Section 5.3 we will see that drawing conclusions about (1) require a sequence of pairwise comparisons similar to those discussed in Section 3.3 and drawing conclusions about (2) above requires regression (linear or logistic). As such, sample size calculations in the context of two-sample t -, Z -, or χ^2 -tests that account for the multiple comparison problem are

appropriate here. We now discuss the analysis approaches that are used to formally draw conclusions about (1) and (2) above.

5.3 Analyzing a Factorial Experiment

As mentioned in the previous section, factorial experiments may be used to find the optimal combination of factor levels, with respect to the response variable of interest, but they may also be used to identify which factors significantly influence the response, either through large main effects or important interaction effects. Determining the optimal condition can be done with pairwise tests, and the overall significance of main effects and interaction effects can be done with regression. The type of tests and the type of regression depends on whether the response variable is continuous or binary. We consider these two scenarios separately in the following subsections.

5.3.1 Continuous Response

We illustrate the topics discussed in this section in the context of the following example. Consider a multi-factor, multi-level version of the Instagram experiment from Section 2.1.2 in which the influence of ad frequency on session duration was evaluated. Now suppose that ad frequency has levels $\{None, 7:1, 4:1, 1:1\}$ corresponding to ad frequencies of never, 1 in 8, 1 in 5, and every other. Additionally, suppose ad type is introduced as a second design factor with levels $\{photo, video\}$. Thus we have two factors with 4 and 2 levels, respectively, thereby creating $2 \times 4 = 8$ experimental conditions – the four ad frequencies with photo ads and the four ad frequencies with video ads.

Now suppose that $n = 1000$ units (users) are randomized to each condition and their session duration is recorded. These data, and the influence of ad frequency and ad type both marginally and jointly are depicted in Figures 11, and 12. In particular, we see in the left plot of Figure 11 that session duration decreases steadily as ad frequency increases, and in the right plot we see (perhaps unsurprisingly) that session duration is slightly longer when video ads are displayed instead of photo ads. In addition to this information, these plots

suggest that ad frequency is a more influential factor than ad type since the change in session duration across different ad frequencies is much larger than the change produced by different ad types. The average session durations for each ad frequency and each ad type are provided in Table 6.




Figure 11: Left: Main effect plot for ad frequency; Right: Main effect plot for ad type.

Table 6: Average session duration for different ad frequencies

Frequency	Average Session Duration (min)
None	7.03
7:1	5.32
4:1	4.41
1:1	1.25
Type	Average Session Duration (min)
Photo	4.25
Video	4.76

However, discussing main effects like this can be uninformative and potentially misleading if there is a significant interaction between the factors. In particular, if the main effect of ad frequency (depicted in the left plot of Figure 11) is different for photo ads versus video ads, then we would say that there is an interaction between the factors. Visually, such an interaction would be indicated if the pattern seen in the left plot of Figure 11 is roughly the same (i.e., parallel lines) for both ad types. Or, equivalently, if the pattern seen in the right plot of Figure 11 is roughly the same (i.e., parallel lines) for all ad frequencies. Figure 12 depicts the interaction plots for the two factors. The left plot corresponds to a plot of

the main effect of frequency for the two different ad types, and the right plot depicts the main effect of ad type for the four different frequencies. Non-parallel lines on these plots would indicate a significant interaction since this would correspond to the main effect of one factor depending on the levels of the other factor. As we can see in Figure 12, the lines are not perfectly parallel, indicating the presence of a small interaction effect. However, the departure from parallelism is not drastic, and so we would not expect the interaction effect to be large. The average session durations in each condition are provided in Table 7

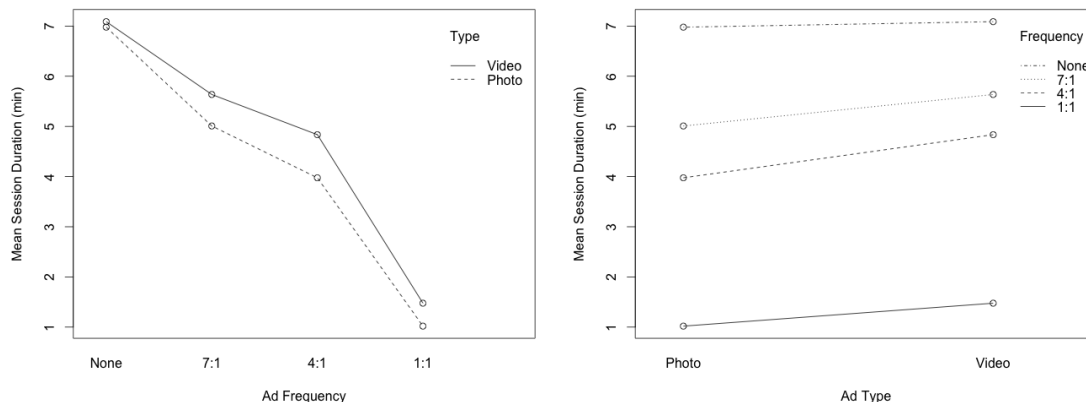


Figure 12: Interaction plot for ad frequency and ad type.

Table 7: Average session duration in each ad frequency-type condition.

		Ad Type	
		Photo	Video
Frequency	None	6.98	7.09
	7:1	5.01	5.64
	4:1	3.98	4.83
	1:1	1.02	1.48

In order to formally evaluate whether the frequency-type main or interaction effects are significant we turn to linear regression. Here we assume that our response variable measurements (in this case, session duration) follow a normal distribution and can be modeled with an ordinary linear regression model, where each term in the model is an indicator (dummy) variable related to the main effect and interaction effects of each design factor. As in Section 3.1.1, a factor with m levels will have $m - 1$ indicator variables in the model corresponding to all but one of the levels. This will be the case for each design factor in the experiment –

their inclusion helps us quantify the main effect of each factor. In addition to this, the model will contain pairwise products of the indicator variables for different factors which constitute the 2-way interaction between the factors, and which help to evaluate the significance of the corresponding interaction effect. Note that when there are more than two design factors, higher order (3-way and 4-way, etc.) interactions may be added into the model, although they are much less interpretable than 2-way interactions.

The linear regression model appropriate for the analysis of this experiment is

$$Y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \beta_3 x_{i3} + \beta_4 x_{i4} + \beta_5 x_{i1} x_{i4} + \beta_6 x_{i2} x_{i4} + \beta_7 x_{i3} x_{i4} + \epsilon_i \quad (5.1)$$

where the x 's are indicator variables such that $x_{i1} = 1$ if unit i is in the 7:1 condition, $x_{i2} = 1$ if unit i is in the 4:1 condition, $x_{i3} = 1$ if unit i is in the 1:1 condition, and $x_{i4} = 1$ if unit i is in the video condition, for $i = 1, 2, \dots, n$.

In the presence of a significant interaction effect, it no longer makes sense to discuss the main effect of a factor, because doing so ignores the fact that this effect changes depending on the level of another factor. As such it is typical to first decide whether the interaction is statistically significant. Notice that if $\beta_5 = \beta_6 = \beta_7 = 0$ then model (5.1) would not account for the interaction between ad frequency and ad type and a formal test of

$$H_0: \beta_5 = \beta_6 = \beta_7 = 0 \text{ vs. } H_A: \beta_j \neq 0$$

for $j = 5, 6, 7$ would evaluate the significance of the interaction effect. If the interaction is significant then any conclusions regarding the effect of one factor must be made in the context of the levels of the other factor.

Alternatively, if the interaction is not significant, then the interaction terms can be removed from the model yielding the following simplified main effects model

$$Y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \beta_3 x_{i3} + \beta_4 x_{i4} + \epsilon_i \quad (5.2)$$

which can be used to evaluate the significance of the main effect of each factor (in this case

Table 8: Expected response in each ad frequency-type condition, based on the main effects model

		Ad Type	
		Photo	Video
Freq.	None	$E[Y_i x_{i1} = x_{i2} = x_{i3} = 0, x_{i4} = 0] = \beta_0$	$E[Y_i x_{i1} = x_{i2} = x_{i3} = 0, x_{i4} = 1] = \beta_0 + \beta_4$
	7:1	$E[Y_i x_{i1} = 1, x_{i4} = 0] = \beta_0 + \beta_1$	$E[Y_i x_{i1} = 1, x_{i4} = 1] = \beta_0 + \beta_1 + \beta_4$
	4:1	$E[Y_i x_{i2} = 1, x_{i4} = 0] = \beta_0 + \beta_2$	$E[Y_i x_{i2} = 1, x_{i4} = 1] = \beta_0 + \beta_2 + \beta_4$
	1:1	$E[Y_i x_{i3} = 1, x_{i4} = 0] = \beta_0 + \beta_3$	$E[Y_i x_{i3} = 1, x_{i4} = 1] = \beta_0 + \beta_3 + \beta_4$

ad frequency and ad type). Table 8 summarizes the expected response (based on this main effects model) in each of the experimental conditions.

Notice that if $\beta_1 = \beta_2 = \beta_3 = 0$ in model there is no difference between the expectations in the four rows. Thus the hypothesis

$$H_0: \beta_1 = \beta_2 = \beta_3 = 0 \text{ vs. } H_A: \beta_j \neq 0$$

for $j = 1, 2, 3$ tests whether ad frequency is a significant factor, and performing this test indicates whether the main effect of ad frequency is significant. Next, notice that if $\beta_4 = 0$ there is no difference between the expectations in the two columns. Thus the hypothesis

$$H_0: \beta_4 = 0 \text{ vs. } H_A: \beta_4 \neq 0$$

tests whether ad type is a significant factor, and performing this test indicates whether the main effect of ad type is significant. It is important to emphasize, however, that these tests and the interpretation of main effects are only appropriate in the absence of interaction.

Each of the null hypotheses just described generates a **reduced model** with fewer terms relative to a **full model** with all terms. Each of the hypotheses, then, is tested by comparing the reduced model to the full one. In particular, to decide whether the frequency-type interaction is significant we compare model (5.1) to the reduced version when $\beta_5 = \beta_6 = \beta_7 = 0$. To decide whether the ad frequency main effect is significant we compare model (5.2) to the reduced version when $\beta_1 = \beta_2 = \beta_3 = 0$, and to decide whether the ad type main effect is significant we compare model (5.2) to the reduced version when $\beta_4 = 0$. Each of these comparisons is done using a **partial F-test** which compares the mean squared

errors between the full and reduced models – similar to the F -test for overall significance in a linear regression. The test statistics and p-values of these tests are provided in standard linear regression ANOVA tables.

The linear regression output from R's `lm()` function applied to these data is shown below, and the ANOVA results from R's `anova()` function are displayed below that. In this output we can see, based on the displayed p-values, that each individual term in the model is statistically significant. However, when judging the overall significance of a main or interaction effect it is inappropriate to draw conclusions from the individual t -tests from which these p-values arise. Instead we must examine the p-values associated with the partial F -tests displayed in the model's ANOVA table.

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	6.97785	0.02824	247.104	< 2e-16 ***
Frequency7:1	-1.96929	0.03994	-49.312	< 2e-16 ***
Frequency4:1	-3.00204	0.03994	-75.173	< 2e-16 ***
Frequency1:1	-5.95856	0.03994	-149.206	< 2e-16 ***
TypeVideo	0.10993	0.03994	2.753	0.00592 **
Frequency7:1:TypeVideo	0.51768	0.05648	9.166	< 2e-16 ***
Frequency4:1:TypeVideo	0.74924	0.05648	13.266	< 2e-16 ***
Frequency1:1:TypeVideo	0.34731	0.05648	6.150	8.14e-10 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.893 on 7992 degrees of freedom

Multiple R-squared: 0.8497, Adjusted R-squared: 0.8496

F-statistic: 6455 on 7 and 7992 DF, p-value: < 2.2e-16

Analysis of Variance Table

Response: Time

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Frequency	3	35353	11784.3	14778.187	< 2.2e-16 ***

Type	1	527	527.3	661.318	< 2.2e-16 ***
Frequency:Type	3	149	49.8	62.398	< 2.2e-16 ***
Residuals	7992	6373	0.8		

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

The p-values in this table are all sufficiently small (smaller than any reasonable significance level), allowing us to conclude that the main effects of ad frequency and ad type, and the interaction between them, are all statistically significant. The size of the test statistics corresponding to these p-values provides insight into the relative size of these effects. In particular we see that ad frequency is substantially more influential than ad type, which is more significant than the interaction effect. Thus we conclude that both factors are important, and so is their interaction, which means that both factors should be considered when trying to optimize session duration.

In order to determine which frequency-type combination (and hence experimental condition) is optimal, a series of one-sided *t*-tests may be used. This is left as an exercise for the reader. In the next section we discuss the analysis of a similarly designed factorial experiment, but where the response variable is binary as opposed to continuous.

5.3.2 Binary Response

The informal and formal evaluation of main and interaction effects can be performed in the context of a binary response variable as well. However, in this situation, logistic regression is used instead of ordinary linear regression. For a primer on logistic regression see Appendix A.4. Like linear regression, the linear predictor in a logistic regression model is composed of dummy variables and products of dummy variables corresponding to the levels of each design factor. A factor with m levels would have $m - 1$ dummy variables in the model, and the highest order interaction term (i.e., the number of dummy variables in the largest product) is k when there are k design factors in the experiment.

The difference between logistic regression and ordinary linear regression is that this linear

predictor is equated to $\text{logit}(\pi_i) = \log(\pi_i/(1 - \pi_i))$. Here $\pi_i = E[Y_i|x_{i1}, x_{i2}, \dots, x_{ip}]$ is the probability that unit i , with explanatory variable values given by $x_{i1}, x_{i2}, \dots, x_{ip}$, performs the action of interest. If the design of such an experiment was equivalent to the Instagram one discussed in the previous section (but with a binary response), the logistic regression model would be

$$\log\left(\frac{\pi_i}{1 - \pi_i}\right) = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \beta_3 x_{i3} + \beta_4 x_{i4} + \beta_5 x_{i1} x_{i4} + \beta_6 x_{i2} x_{i4} + \beta_7 x_{i3} x_{i4} \quad (5.3)$$

where the x 's are indicator variables such that $x_{i1} = 1$ if unit i is in the 7:1 condition, $x_{i2} = 1$ if unit i is in the 4:1 condition, $x_{i3} = 1$ if unit i is in the 1:1 condition, and $x_{i4} = 1$ if unit i is in the video condition, for $i = 1, 2, \dots, n$.

As in the case of linear regression, interest lies in determining whether subsets of the β 's are equal to zero to evaluate the significance of various main and interaction effects. Whereas partial F -tests are used for this task in the context of linear regression models, we use **likelihood ratio tests** for this purpose with logistic regression models. For instance, a test of

$$H_0: \beta_5 = \beta_6 = \beta_7 = 0 \text{ vs. } H_A: \beta_j \neq 0$$

for $j = 5, 6, 7$ would evaluate the significance of the interaction effect. In order to perform such a test, we compare the maximized log-likelihood of the full model (5.3) to that of the reduced model (when H_0 is true). To avoid a technical definition we simply state that log-likelihood of a given model is a measure of how well the model fits the data (i.e., how closely what was observed agrees with what the model expects). The larger the log-likelihood value, the better the corresponding model fits the data.

The test statistic for the the likelihood ratio test is referred to as the **deviance**, denoted Λ , and is defined as

$$\begin{aligned} \Lambda &= 2[\log\text{-likelihood}(\text{full model}) - \log\text{-likelihood}(\text{reduced model})] \\ &= 2\log\left(\frac{\text{likelihood}(\text{full model})}{\text{likelihood}(\text{reduced model})}\right) \end{aligned}$$

which, if H_0 is true, approximately follows a χ^2 distribution with l degrees of freedom,

where l is the number of restrictions implied by H_0 . For instance, the null hypothesis $H_0: \beta_5 = \beta_6 = \beta_7 = 0$ places restrictions on three of the β 's and so $l = 3$ in this case. A simple way to determine the value of l is to compare the number of regression coefficients in the two models: if the number of β 's in the full model is p , and $q < p$ in the reduced model, then $l = p - q$. The p-value for this test is calculated as $\text{p-value} = P(T \geq \Lambda)$ where $T \sim \chi^2_{(l)}$.

As noted, such tests are useful for identifying significant main and interaction effects. To determine which of the experimental conditions is optimal a series of one-sided pairwise Z - or χ^2 -tests may be performed. In Section 6.1 we illustrate this logistic regression approach to analyzing factorial experiments with binary response variables.

Two-Level Factorial Experiments

In the previous chapter we introduced factorial experiments as an informative and efficient way of investigating the influence of multiple factors on some response variable of interest. The key advantage to this method of experimentation (relative to the one-factor-at-a-time approach) is that every possible combination of factor levels is considered, and so we do not risk missing an optimal combination. However, we saw that this advantage is also the main drawback of such an experiment. The problem is that as the number of factors and the number of levels increase, the number of unique combinations, and hence experimental conditions, gets very large.

In this chapter we consider a special case of factorial experiments: **two-level** factorial experiments. In these experiments we consider investigating k design factors, each at only two levels. While several levels may be plausible, just two must be chosen for the purpose of these experiments. Such experiments are typically used for factor **screening**. When a large number of factors may potentially influence the response, screening experiments may be used to determine which of them is most influential. In practice, the **Pareto principle**³ often applies and interest lies in determining the small number of important factors. Once these factors have been identified, follow-up experiments that investigate just these factors at a larger number of levels can be performed. Here we think of these two-level experiments as providing an efficient method of homing in on truly influential factors.

With k factors, each at two levels, there are total of 2^k unique combinations of the factors' levels. In Section 6.1 we consider **2^k factorial experiments** which investigate k factors with exactly 2^k experimental conditions. In other words, the 2^k factorial experiment considers every possible combination of the factors's levels. In Section 6.2 we discuss **2^{k-p}**

³The Pareto principle states that only a *vital few* factors are important relative to the *trivial many*.

fractional factorial experiments which investigate k factors with just 2^{k-p} specifically chosen experimental conditions (i.e., just a *fraction* of all possible conditions).

6.1 2^k Factorial Experiments

6.1.1 Designing 2^k Factorial Experiments

The design of a 2^k factorial experiment involves choosing k factors that are expected to influence the response variable (Y) in some way. Then, two levels for each factor are chosen to experiment with. Given that this form of factorial experiment is used for screening, and hence identifying influential factors, it is important to choose levels that provide the largest opportunity for an influential factor to be noticed. Recall that the main effect of a factor is defined to be the change in response produced by a change in the factor; thus levels should be chosen that are quite different from one another; even a very influential factor may not appear to be influential if the factor levels are too similar. In the context of two-level experiments, we arbitrarily think of these levels as *low* and *high* values of the factors.

With the factors and factor levels chosen, the experimental conditions are defined to be the unique combinations of the factor levels. Since each factor is investigated at just two levels, there are 2^k distinct combinations and hence 2^k experimental conditions. Experimental units are then assigned to each condition. For ease of notation, we assume that the experiment is balanced and n units are assigned to each of the 2^k conditions. The choice of sample size n can be determined using power analyses based on two-sample tests that account for the multiple comparison problem, as was discussed in Section 5.2. Note that the balanced sample size assumption is not necessary; the analysis techniques discussed in the next section are applicable even when sample sizes are not balanced.

6.1.2 Analyzing 2^k Factorial Experiments

The analysis of a 2^k factorial experiment is carried out via regression: linear regression (when Y is continuous) and logistic regression (when Y is binary). The corresponding fitted

models provide an estimate of the **response surface** that relates the response variable to the k factors. The linear predictor for these models contain k main effect terms, $\binom{k}{2}$ two-factor interaction terms, $\binom{k}{3}$ three-factor interaction terms and so on, up to and including $\binom{k}{k} = 1$ k -factor interaction term. For instance, the linear predictor associated with a 2^3 factorial experiment looks as follows

$$\beta_0 + \beta_1x_1 + \beta_2x_2 + \beta_3x_3 + \beta_{12}x_1x_2 + \beta_{13}x_1x_3 + \beta_{23}x_2x_3 + \beta_{123}x_1x_2x_3.$$

Whether the linear predictor is a part of a linear or logistic regression, we code the levels of the x 's as -1 and $+1$, corresponding to the *low* and *high* levels of each factor. When the factor is categorical this coding is simply assigned to the factor's levels. But when the factor is numeric the coding arises through the following transformation

$$x_C = \frac{x_N - (x_H + x_L)/2}{(x_H - x_L)/2}$$

where x_H and x_L correspond to the high and low values of the factor as recorded in the natural units, and x_C corresponds to the coded version of x_N (a particular value of the factor in the natural units). It can be readily seen that substituting $x_N = x_H$ and $x_N = x_L$ into this equation gives $+1$ and -1 , respectively.

Although it is possible to treat a numeric factor as categorical and simply assign the ± 1 coding to the factor's levels, the relationship in the preceding equation facilitates the conversion back and forth between the coded and natural units. This conversion is useful when interpolation is required; suppose your response surface suggests that the optimal level of a numeric factor is 0.75 on the coded scale – what does this value translate to in the natural units? The conversion formula above can be used to determine this.

With the ± 1 coding in place, each experimental condition can be identified by a unique combination of plus and minus ones. In fact, the design of the experiment can be displayed in what is known as a **design matrix**. The design matrix associated with a 2^3 factorial design is shown in Table 9; we see that each row corresponds to a unique condition and the columns correspond to factors which, for a given condition, indicate which level the factor

is to be set at. Thus this table provides a prescription for running the 2^3 experiment. The design matrix for a general 2^k experiment looks similar – there are simply k columns, each with 2^{k-1} -1 's and 2^{k-1} $+1$'s strategically alternated so that each row contains a unique combination of ± 1 's.

Table 9: The design matrix for a 2^3 factorial experiment.

Condition	Factor 1	Factor 2	Factor 3
1	-1	-1	-1
2	+1	-1	-1
3	-1	+1	-1
4	+1	+1	-1
5	-1	-1	+1
6	+1	-1	+1
7	-1	+1	+1
8	+1	+1	+1

Then, using the data, which is collected in a similar manner (i.e., factor levels recorded as ± 1 's), we fit a linear or logistic regression model. As noted above, this choice is dictated by the type of response variable (continuous vs. binary). Apart from this difference, the models themselves are based on the exact same linear predictor and investigation of the significance of main and interaction effects can be carried out by deciding whether particular β 's are significantly different from zero. However, the particular tests that are used for this differ depending on the type of regression model. In the case of linear regression t -tests are used to evaluate the significance of individual β 's while Z -tests are used for this purpose in the context of logistic regression. To decide whether several β 's are simultaneously zero (and hence compare full and reduced models), partial F -tests are used in the case of linear regression, and likelihood ratio tests are used in the case of logistic regression.

To illustrate the analysis of such an experiment, we consider an example from [Montgomery \(2017\)](#) in which an experiment was performed to test new ideas to improve the conversion rate of credit card offers. For this example, the response is binary – indicating whether an individual signed up for a credit card as a result of the offer – and so an analysis based on logistic regression is performed. An illustration of the linear regression analysis of a two-level factorial experiment is shown in [Section 6.2.2](#).

A 2^4 factorial experiment was carried out to investigate four factors and their influence on credit card sign ups. The four factors and each of their levels are summarized in Table 10. The $2^4 = 16$ unique combinations of these factor levels produced 16 experimental conditions, each of which was assigned $n = 7500$ units. Practically speaking, 16 credit card offers were devised (one corresponding to each condition) and each was mailed to 7500 customers. The design matrix and a summary of the conversion rates are provided in Table 11.

Table 10: Factors and levels for the credit card example.

Factor	Low (-)	High (+)
Annual Fee (x_1)	Current	Lower
Account-Opening Fee (x_2)	No	Yes
Initial Interest Rate (x_3)	Current	Lower
Long-term Interest Rate (x_4)	Low	High

Table 11: Design matrix and response summary for the 2^4 factorial credit card experiment.

Condition	Factor 1	Factor 2	Factor 3	Factor 4	Sign-ups	Conversion Rate
1	-1	-1	-1	-1	184	2.45%
2	+1	-1	-1	-1	252	3.36%
3	-1	+1	-1	-1	162	2.16%
4	+1	+1	-1	-1	172	2.29%
5	-1	-1	+1	-1	187	2.49%
6	+1	-1	+1	-1	254	3.39%
7	-1	+1	+1	-1	174	2.32%
8	+1	+1	+1	-1	183	2.44%
9	-1	-1	-1	+1	138	1.84%
10	+1	-1	-1	+1	168	2.24%
11	-1	+1	-1	+1	127	1.69%
12	+1	+1	-1	+1	140	1.87%
13	-1	-1	+1	+1	172	2.29%
14	+1	-1	+1	+1	219	2.92%
15	-1	+1	+1	+1	153	2.04%
16	+1	+1	+1	+1	152	2.03%

Using this data we fit a logistic regression model with the following linear predictor

$$\begin{aligned}
&\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + \beta_{12} x_1 x_2 + \beta_{13} x_1 x_3 + \beta_{14} x_1 x_4 + \\
&\beta_{23} x_2 x_3 + \beta_{24} x_2 x_4 + \beta_{34} x_3 x_4 + \beta_{123} x_1 x_2 x_3 + \beta_{124} x_1 x_2 x_4 + \\
&\beta_{134} x_1 x_3 x_4 + \beta_{234} x_2 x_3 x_4 + \beta_{1234} x_1 x_2 x_3 x_4
\end{aligned} \tag{6.1}$$

We do this in R using the `glm()` function with a logit link function. A summary of this model fit is provided below.

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-3.739697	0.019342	-193.347	< 2e-16	***
x1	0.080845	0.019342	4.180	2.92e-05	***
x2	-0.106211	0.019342	-5.491	3.99e-08	***
x3	0.058248	0.019342	3.011	0.00260	**
x4	-0.108086	0.019342	-5.588	2.29e-08	***
x1:x2	-0.055164	0.019342	-2.852	0.00434	**
x1:x3	-0.004794	0.019342	-0.248	0.80426	
x2:x3	-0.006967	0.019342	-0.360	0.71868	
x1:x4	-0.013178	0.019342	-0.681	0.49566	
x2:x4	0.010625	0.019342	0.549	0.58280	
x3:x4	0.038079	0.019342	1.969	0.04899	*
x1:x2:x3	-0.009646	0.019342	-0.499	0.61799	
x1:x2:x4	0.010629	0.019342	0.550	0.58265	
x1:x3:x4	-0.002543	0.019342	-0.131	0.89539	
x2:x3:x4	-0.020946	0.019342	-1.083	0.27885	
x1:x2:x3:x4	-0.009496	0.019342	-0.491	0.62347	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 26854 on 119999 degrees of freedom
 Residual deviance: 26741 on 119984 degrees of freedom
 AIC: 26773

Number of Fisher Scoring iterations: 6

As we can see, all main effects are significant but there are several insignificant terms; in particular, all three-factor and four-factor interactions are insignificant, and only two of the two-factor interactions are significant. In light of this, we fit a reduced model with $\beta_{13} = \beta_{14} = \beta_{23} = \beta_{24} = \beta_{123} = \beta_{124} = \beta_{134} = \beta_{234} = \beta_{1234} = 0$. The output from this reduced model is shown below.

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-3.73961	0.01934	-193.316	< 2e-16 ***
x1	0.08214	0.01920	4.279	1.88e-05 ***
x2	-0.10834	0.01920	-5.644	1.66e-08 ***
x3	0.05886	0.01916	3.072	0.00212 **
x4	-0.11068	0.01916	-5.777	7.61e-09 ***
x1:x2	-0.05706	0.01920	-2.972	0.00296 **
x3:x4	0.04051	0.01916	2.115	0.03447 *

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 26854 on 119999 degrees of freedom
Residual deviance: 26744 on 119993 degrees of freedom
AIC: 26758

Number of Fisher Scoring iterations: 6

While the Z -test p-values in the full model suggested this reduced model is appropriate, it is wise to formally test

$$H_0 : \beta_{13} = \beta_{14} = \beta_{23} = \beta_{24} = \beta_{123} = \beta_{124} = \beta_{134} = \beta_{234} = \beta_{1234} = 0$$

with a likelihood ratio test as described in Section 5.3.2. The maximized log-likelihood

for the full model is -13370.45 and for the reduced model is -13371.91 (obtained using the `logLik()` function in R) which gives a deviance statistic of $\Lambda = 2.9244$ that corresponds to a p-value of $P(T \geq 2.9244) = 0.9672$, where $T \sim \chi^2_{(9)}$. Thus we do not reject H_0 , implying that the reduced model is adequate and that only the main effects and two two-factor interactions are significant. Main effect and interaction effect plots are shown in Figures 13 and 14. Unsurprisingly, the main effect plots demonstrate that lower annual fees, no account-opening fee, lower initial interest rates and lower long-term interest rates are all associated with increased conversion rates. Furthermore, the interaction plots indicate that a lower annual fee is especially effective at increasing the conversion rate when there is no account-opening fee; when there is an account opening fee, the annual fee is not as influential. We can also see that as long as the long-term interest rate is low, the initial interest rate doesn't really matter.




Figure 13: Main effect plots for the credit card example.

These plots and the regression results provide very useful insight into determining which factors most influence credit card intitiation, and which combination of factor levels is op-




Figure 14: Interaction effect plots for the credit card example.

timal. In the context of these 16 experimental conditions, the three best conversion rates occurred in conditions 2, 6 and 14 where the conversion rates were 3.36%, 3.39% and 2.92%, respectively. The p-value associated with a χ^2 -test of $H_0 : \pi_2 = \pi_6 = \pi_{14}$ is 0.1917, indicating no significant difference between these conditions. Presumably the credit card company would choose to implement the one that is most profitable for them.

6.2 2^{k-p} Fractional Factorial Experiments

As we have seen, the 2^k factorial experiment is a useful special case of a general factorial experiment since it minimizes the number of levels being investigated, and hence reduces the overall number of experimental conditions. However, the 2^k factorial experiment still investigates *all possible* combinations of the factor levels, so an experiment with just $k = 8$ factors requires $2^8 = 256$ distinct experimental conditions – which seems unmanageably large.

A 2^{k-p} fractional factorial experiment, on the other hand, similarly investigates k factors but in just a fraction of the conditions: $(1/2)^p$ to be exact. These conditions are specially chosen to ensure that all main effects and any potentially important interaction effects can be estimated. As we will see, not *all* interaction effects will be estimated, but this is the

sacrifice we make in order to investigate a relatively large number of factors with a relatively small number of conditions. We elaborate on this point in the next two subsections.

6.2.1 Designing 2^{k-p} Fractional Factorial Experiments

To motivate the utility and the mechanics of a fractional factorial experiment, consider the linear predictor from the full 2^k factorial experiment, which contains k main effect terms, $\binom{k}{2}$ two-factor interaction terms, $\binom{k}{3}$ three-factor interaction terms and so on, up to and including $\binom{k}{k} = 1$ k -factor interaction term; this is a total of $\sum_{i=1}^k \binom{k}{i} = 2^k - 1$ regression coefficients. However, just $k + \binom{k}{2}$ of these are main effects and two-factor interactions – the remaining correspond to higher order interaction terms. In the case that $k = 8$, there are 8 main effects, 28 two-factor interactions and 219 higher order interactions, many of which are likely to be insignificant.

The principle of **effect sparsity** says that in the presence of several factors, variation in the response is likely to be driven by a small number of main effects and low-order interactions. In fact we saw this in the example in Section 6.1.2: none of the three-factor or four-factor interactions were significant, and in fact many of the two-factor interactions were insignificant. As such, it is typically a waste of resources to estimate these higher order interaction terms, and a better use of these resources is to estimate the main effects and low-order interactions of a larger number of factors.

In order to do this we must first explain what p is in the 2^{k-p} notation. Whereas with a full factorial approach, investigating k factors with two levels each would require 2^k conditions, we endeavor to do this in a fraction of the conditions; the size of this fraction is determined by p . If we'd like to investigate k factors with half as many conditions then we use a 2^{k-1} experiment. If we'd like to investigate k factors with just a quarter of the conditions then we use a 2^{k-2} experiment. In general, if we'd like to investigate k factors in $(1/2)^p$ as many conditions we use a 2^{k-p} experiment.

Then, if a full factorial approach requires 2^k conditions and we only want 2^{k-p} , we need to choose *which* 2^{k-p} conditions to experiment with. For instance, if $k = 5$ and $p = 2$,

then the goal is to investigate 5 factors in $2^3 = 8$ conditions (where normally 32 conditions would be required with the full factorial approach). The design matrix for a full 2^5 factorial experiment is shown in Table 12. The question is, among the 32 conditions shown there, which 8 do we choose for the 2^{5-2} fractional design?

Table 12: The design matrix for a full 2^5 factorial experiment.

Condition	Factor 1	Factor 2	Factor 3	Factor 4	Factor 5
1	-1	-1	-1	-1	-1
2	+1	-1	-1	-1	-1
3	-1	+1	-1	-1	-1
4	+1	+1	-1	-1	-1
5	-1	-1	+1	-1	-1
6	+1	-1	+1	-1	-1
7	-1	+1	+1	-1	-1
8	+1	+1	+1	-1	-1
9	-1	-1	-1	+1	-1
10	+1	-1	-1	+1	-1
11	-1	+1	-1	+1	-1
12	+1	+1	-1	+1	-1
13	-1	-1	+1	+1	-1
14	+1	-1	+1	+1	-1
15	-1	+1	+1	+1	-1
16	+1	+1	+1	+1	-1
17	-1	-1	-1	-1	+1
18	+1	-1	-1	-1	+1
19	-1	+1	-1	-1	+1
20	+1	+1	-1	-1	+1
21	-1	-1	+1	-1	+1
22	+1	-1	+1	-1	+1
23	-1	+1	+1	-1	+1
24	+1	+1	+1	-1	+1
25	-1	-1	-1	+1	+1
26	+1	-1	-1	+1	+1
27	-1	+1	-1	+1	+1
28	+1	+1	-1	+1	+1
29	-1	-1	+1	+1	+1
30	+1	-1	+1	+1	+1
31	-1	+1	+1	+1	+1
32	+1	+1	+1	+1	+1

To answer this, we consider the design matrix associated with a full 2^3 factorial experi-

ment. More generally, we consider the design matrix of a full 2^{k-p} factorial experiment (i.e., an experiment that investigates $k - p$ factors with 2^{k-p} conditions). In Table 13 we consider an extended version of the 2^3 design matrix that includes columns corresponding the two- and three-factor interactions. These columns are obtained by performing elementwise products of the columns titled A, B and C. For example, the entries in the AB interaction column are obtained by multiplying the corresponding entries of columns A and B. Note that for brevity, we denote factors 1, 2 and 3 here with A, B and C. Remember that each of these columns correspond to values of the factors, and so when it comes to fitting a regression model each column is used to estimate the corresponding effect. Although a full justification is outside the scope of these notes, we mention that because we've coded factor levels as ± 1 , each of these columns is orthogonal to one another; as a consequence each column is individually responsible for estimating a single effect. Thus, the AB column is used to estimate β_{AB} , the interaction effect between factors A and B. Similarly, the ABC column is used to estimate β_{ABC} , the three-way interaction between factors A, B and C.

Table 13: An extended design matrix for the 2^3 factorial experiment.

Condition	A	B	C	AB	AC	BC	ABC
1	-1	-1	-1	+1	+1	+1	-1
2	+1	-1	-1	-1	-1	+1	+1
3	-1	+1	-1	-1	+1	-1	+1
4	+1	+1	-1	+1	-1	-1	-1
5	-1	-1	+1	+1	-1	-1	+1
6	+1	-1	+1	-1	+1	-1	-1
7	-1	+1	+1	-1	-1	+1	-1
8	+1	+1	+1	+1	+1	+1	+1

The relevance of this becomes apparent in the context of the effect sparsity principle. Because interaction effects are commonly negligible (especially high order interactions), we use an interaction column (that is likely to be negligible) to dictate when to run a new factor, say D, at its low and high levels. For instance, imagine in addition to factors A, B and C, we wish to consider factor D. We can do this by taking (for example) the ABC column and using its plus and minus ones as a prescription indicating the level of D for a given experimental condition. In this case we see that $ABC = +1$ in condition 3 and so factor D should be run at its high level in condition 3. Similarly, in condition 7 $ABC = -1$ and so factor D would

be run at its low level. Recall that the motivating example had $k = 5$ factors which were to be investigated with 8 conditions. In this situation we'd have a fifth factor E. To determine when E should be run at its low and high levels we similarly associate it with one of the interaction columns in Table 13. Take BC, for example; the plus and minus ones of the BC column now indicate factor E's levels in each condition.

This association of new factors to existing interactions is referred to as **aliasing**. Previously we aliased D with the ABC interaction and E with the BC interaction. We denote this aliasing as $D=ABC$ and $E=BC$ and call ' $D=ABC$ ' and ' $E=BC$ ' the **design generators**. When we do this, we **confound** the interaction effect with the main effect of the new factor (i.e., ABC with D, or BC with E). For example, in a full factorial experiment the ABC column in Table 13 estimates the ABC interaction. But once aliased in the context of a fractional factorial experiment the ABC column *also* estimates the main effect of D. Thus β_{ABC} now quantifies the joint effects of the ABC interaction and the main effect of factor D. In this case we say that their effects are confounded – that is, we cannot separately estimate the main effect of D from the ABC interaction effect.

Due to the confounding that results from aliasing a new main effect with an existing interaction, it is important to think carefully about *which* interaction to choose as an alias. If at all possible, it is best to avoid aliasing a new factor with an interaction that is likely to be significant since separately estimating these significant effects is desirable. As such, high order interaction terms (that are unlikely to be significant) are good choices for aliases. This notion is quantified by the **resolution** of the fractional factorial design. In general, a design is of resolution R if main effects are aliased with interaction effects involving at least $R - 1$ factors. Resolution is typically denoted with roman numerals. In the fractional factorial design we've been discussing, main effects are aliased with two- and three-factor interactions, the smaller of these numbers being 2. Thus $R - 1 = 2$ and the associated experiment is resolution III.

In general, higher resolution designs are to be preferred over lower resolution designs. For instance, resolution IV and V designs are to be preferred over a resolution III design since in these cases main effects will not be confounded with two-factor interactions. Since

two-factor interactions are typically important, it is best if their effects are not confounded with main effects. The resolution of a fractional factorial experiment is determined by two things:

1. The degree of fractionation desired (i.e., the size of p relative to k).
2. The design generators chosen for aliasing.

The degree of fractionation is typically determined by resource constraints; i.e., how many experimental conditions can you feasibly manage? Given the degree of fractionation (p) we typically choose design generators to maximize resolution. Note that the number of design generators is equal to p . For an excellent guide that simplifies this decision, see the table on page 342 of [Montgomery \(2017\)](#) which, for given values of k and p provides the design generators necessary to achieve a particular resolution. Recall that our 2^{5-2} design was resolution III (denoted 2_{III}^{5-2}). If we wished to investigate $k = 5$ factors in a fractional factorial experiment with a better (higher) resolution, this table shows us that we would need 16 conditions. In particular, the 2^{5-1} design has resolution V (denoted 2_V^{5-1}).

In the next section we illustrate the analysis of a 2_{IV}^{8-4} fractional factorial experiment which investigates $k = 8$ factors in 16 conditions.

6.2.2 Analyzing 2^{k-p} Fractional Factorial Experiments

The analysis of a 2^{k-p} fractional factorial experiment is based on regression models (linear or logistic, depending on the response variable) and in fact is not very different from the analysis procedures for 2^k factorial experiments discussed in Section 6.1.2. Because a 2^{k-p} fractional factorial experiment is performed with only 2^{k-p} conditions, the regression equations we are able to fit are the same ones as in a full 2^k factorial design. The wrinkle in this situation is that the effects estimated in these models are confounded with other effects, and so we are not able to say with 100% certainty whether a given effect is due to say a main effect, or perhaps the interaction it is aliased with. However, if the design has a sufficiently large resolution, we hope that important effects are aliased with high-order

interactions (that are likely negligible) providing us with some confidence that significant effects are not due to the high-order interactions.

To make these ideas more clear we consider an example from [Montgomery \(2017\)](#) in which a fractional factorial experiment was used in the production of wine to study the influence of a variety of factors on a particular vintage of Pinot Noir. In this experiment $k = 8$ factors were investigated each at two levels (the factors and their levels are shown in Table 14) which, if a full factorial experiment was used, would have required 256 conditions. However a 2_{IV}^{8-4} fractional factorial experiment was performed that required only 16 conditions. The response variable in this case is the rating of the wine as determined by 5 raters.

Table 14: Factors and levels for the wine example.

Factor	Low (-)	High (+)
Pinot Noir clone (A)	Pommard	Wadenswil
Oak type (B)	Allier	Troncias
Age of barrel (C)	Old	New
Yeast/skin contact (D)	Champagne	Montrachet
Stems (E)	None	All
Barrel toast (F)	Light	Medium
Whole cluster (G)	None	10%
Fermentation temperature (H)	Low (75°F max)	High (92°F max)

Thus, 16 different wines were produced (based on the 16 different experimental conditions) and $n = 5$ raters tasted and rated each of them (low scores are good, large scores are bad). The design matrix and a summary of the response is provided in Table 15. It is easily verified that the $p = 4$ design generators are $E=BCD$, $F=ACD$, $G=ABC$ and $H=ABD$.

Because the response variable in this setting is continuous, we use linear regression to analyze the data from this experiment. And because only $2^4 = 16$ conditions were used, we can only fit a model with 16 regression coefficients. In the context of a full 2^4 factorial experiment, this would be the model with 4 main effects, 6 two-factor interactions, 4 three-factor interactions and 1 four-factor interaction (like the one shown in equation (6.1)). A summary of this fitted model is shown below.

Coefficients:

Table 15: Design matrix and response summary for the 2^{8-4} fractional factorial wine experiment.

Condition	A	B	C	D	E	F	G	H	Average Rating
1	-1	-1	-1	-1	-1	-1	-1	-1	9.6
2	+1	-1	-1	-1	-1	+1	+1	+1	10.8
3	-1	+1	-1	-1	+1	-1	+1	+1	12.6
4	+1	+1	-1	-1	+1	+1	-1	-1	9.2
5	-1	-1	+1	-1	+1	+1	+1	-1	9.0
6	+1	-1	+1	-1	+1	-1	-1	+1	15.0
7	-1	+1	+1	-1	-1	+1	-1	+1	5.0
8	+1	+1	+1	-1	-1	-1	+1	-1	15.2
9	-1	-1	-1	+1	+1	+1	-1	+1	2.2
10	+1	-1	-1	+1	+1	-1	+1	-1	7.0
11	-1	+1	-1	+1	-1	+1	+1	-1	8.8
12	+1	+1	-1	+1	-1	-1	-1	+1	2.8
13	-1	-1	+1	+1	-1	-1	+1	+1	4.6
14	+1	-1	+1	+1	-1	+1	-1	-1	2.4
15	-1	+1	+1	+1	+1	-1	-1	-1	9.2
16	+1	+1	+1	+1	+1	+1	+1	+1	12.6

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	8.5000	0.2658	31.985	< 2e-16 ***
A	0.8750	0.2658	3.293	0.001619 **
B	0.9250	0.2658	3.481	0.000906 ***
C	0.6250	0.2658	2.352	0.021772 *
D	-2.3000	0.2658	-8.655	2.27e-12 ***
A:B	-0.3500	0.2658	-1.317	0.192532
A:C	1.3000	0.2658	4.892	7.07e-06 ***
B:C	0.4500	0.2658	1.693	0.095261 .
A:D	-0.8750	0.2658	-3.293	0.001619 **
B:D	1.2250	0.2658	4.610	1.98e-05 ***
C:D	0.3750	0.2658	1.411	0.163063
A:B:C	1.5750	0.2658	5.927	1.35e-07 ***
A:B:D	-0.3000	0.2658	-1.129	0.263168
A:C:D	-1.0000	0.2658	-3.763	0.000367 ***
B:C:D	1.1000	0.2658	4.139	0.000104 ***

```
A:B:C:D      0.4750      0.2658      1.787 0.078613 .
```

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 2.377 on 64 degrees of freedom
```

```
Multiple R-squared:  0.7873, Adjusted R-squared:  0.7374
```

```
F-statistic: 15.79 on 15 and 64 DF,  p-value: 4.547e-16
```

Notice this output does not involve the factors E, F, G or H – it only directly references factors A, B, C and D. However, because of the confounding associated with the aliasing in this experiment the BCD interaction estimate also corresponds to the main effect of E, the ACD interaction estimate also corresponds to the main effect of F, the ABC interaction estimate also corresponds to the main effect of G, and the ABD interaction estimate also corresponds to the main effect of H. While we cannot technically separate these effects, we assume that the three-factor interactions are negligible, and hence any significant effect observed is due to the aliased main effect. The same model summary is shown again below, but this time with factors E, F, G and H referenced instead of the three-factor interactions.

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	8.5000	0.2658	31.985	< 2e-16 ***
A	0.8750	0.2658	3.293	0.001619 **
B	0.9250	0.2658	3.481	0.000906 ***
C	0.6250	0.2658	2.352	0.021772 *
D	-2.3000	0.2658	-8.655	2.27e-12 ***
E	1.1000	0.2658	4.139	0.000104 ***
F	-1.0000	0.2658	-3.763	0.000367 ***
G	1.5750	0.2658	5.927	1.35e-07 ***
H	-0.3000	0.2658	-1.129	0.263168
A:B	-0.3500	0.2658	-1.317	0.192532
A:C	1.3000	0.2658	4.892	7.07e-06 ***

A:D	-0.8750	0.2658	-3.293	0.001619	**
A:E	0.4750	0.2658	1.787	0.078613	.
A:F	0.3750	0.2658	1.411	0.163063	
A:G	0.4500	0.2658	1.693	0.095261	.
A:H	1.2250	0.2658	4.610	1.98e-05	***

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

Residual standard error: 2.377 on 64 degrees of freedom

Multiple R-squared: 0.7873, Adjusted R-squared: 0.7374

F-statistic: 15.79 on 15 and 64 DF, p-value: 4.547e-16

Based on this output it appears as though all of the main effects – except H (fermentation temperature) – are significant, though factors D, E, F, G (yeast/skin contact, stems, barrel toast, whole cluster) are most influential. Additionally, the AC, AH and AD interactions also appear to be significant. Note that because of the aliasing structure imposed by the four design generators, these two-factor interactions are respectively confounded with DF, FG and EG – and because factors D, E, F and G are most influential, it is likely that the DF, FG and EG interactions are responsible for the significant effect – not the AC, AH and AD interactions. However, this is simply speculation – we cannot know for certain what causes a significant effect – this is the sacrifice that is made when performing a fractional factorial experiment.

Note that a partial F -test of

$$H_0 : \beta_H = \beta_{AB} = \beta_{AE} = \beta_{AF} = \beta_{AG} = 0$$

which compares the full model above to the one that is reduced by H_0 has an associated p-value of $P(T \geq 2.2124) = 0.06375$ where $T \sim F(5, 64)$. Thus we do not reject H_0 and we conclude that all factors other than H are significantly influential, and the DF, FG and EG interactions are statistically significant. Figures 15 and 16 depict main effect plots for all eight factors and Figure 17 depicts the interaction effect plots for the three significant

interactions.




Figure 15: Main effect plots for the wine example.

By examining the main effect plots it becomes clear that yeast type (D) and the amount of whole clusters (G) used during fermentation are most important, with no whole clusters and Montrachet yeast producing a better tasting Pinot Noir. As well, medium barrel toast (F) and no stems (E) also seem to correspond to a better tasting wine. The interaction plots indicate that if yeast type is Montrachet, the level of barrel toasting doesn't matter much, but if yeast type is Champagne, a medium barrel toast is best. Also, if barrel toast is chosen to be medium, then not including any whole-clusters is best, and similarly, if using none of the stems, then it is also best not to include any whole-clusters.




Figure 16: More main effect plots for the wine example.




Figure 17: Interaction effect plots for the wine example.

Appendix

In this Appendix we review some of the statistical prerequisites for the material discussed throughout the notes. In particular we review random variables and probability distributions, point and interval estimation, hypothesis testing, linear regression, and logistic regression.

A.1 Random Variables and Probability Distributions

A.1.1 Random Variables and Probability Functions

A **random variable** $Y : \Omega \rightarrow \mathbb{R}$ is a function that assigns real numbers to outcomes of a random process, such as flipping a coin or measuring some quantity of interest. We refer to the possible values a random variable can take on as the **support set**, and we dichotomize random variables based on the type of values they assume. A **discrete** random variable is one whose support set is finite or countably infinite such as $y = 0, 1, 2, \dots, n$ or $y = 0, 1, 2, \dots$. We typically use discrete random variables when counting events is of interest. A **continuous** random variable, on the other hand, takes on a continuum of values and so its support set is a subinterval of the real numbers such as $y \geq 0$, $y \in [0, 1]$ or $-\infty < y < \infty$. We typically use continuous random variables when measuring some continuous quantity is of interest. Note that for clarity we denote random variables with upper case letters and the values they take on with lower case letters.

Example 1: Suppose we send an email survey to $n = 30$ individuals and we're interested in the the number of these individuals that respond to the survey. Let Y represent the number of survey responses. In this case the support set is $y = 0, 1, 2, \dots, 30$, and so Y is a discrete random variable.

Example 2: Interest often lies in measuring lifetimes of people, products, and processes. Suppose that, in particular, we are interested in the lifetime of an iPhone’s battery. Let Y represent the lifetime (in hours) of an iPhone battery. In this case the support set is theoretically $y \geq 0$, which is a continuous subinterval of the real numbers, and so Y is a continuous random variable.

Because random variables take on values randomly, interest lies in quantifying the probability that Y assumes a particular value (i.e., $P(Y = a)$) or lies in some interval (i.e., $P(a < Y < b)$). Such probabilities are described by the **probability distribution** of the random variable and quantified by the corresponding **probability function** $f(y)$. The form of this function will differ from one distribution to another, but in all cases, by substituting all values of $y \in A$ (where A is the support set of Y) into $f(y)$ and constructing a plot of $f(y)$ vs. y , we can visualize the probability distribution. Doing so provides insight into the shape of the distribution – most notably, the center and spread – and hence an idea of what values of y seem typical and which ones seem extreme.




Figure A.1: Left: Distribution Characterizing Survey Respondents; Right: Distribution Characterizing iPhone Battery Lifetimes

Figure A.1 depicts hypothetical distributions for the random variables defined in Examples 1 (left panel) and 2 (right panel). We see that when Y is a discrete random variable the plot of $f(y)$ vs. y is a barplot, with bar heights equaling the probability the Y takes on a given value y . On the other hand, the plot of $f(y)$ vs. y for continuous Y is a smooth curve.

In the left hand plot we see that one could reasonably expect 0 to 15 survey responses, with 4 to 8 responses being most likely, and anymore than 15 responses very unlikely. Similarly, the right plot suggests that it is quite likely that an iPhone will last up to 10 hours on a single charge, but it is not very likely to live past 20 hours on a single charge.

To formalize observations like these, we can use probability functions to calculate the probability that such events occur. However, the manner in which these functions are used to calculate probabilities depends on whether Y is discrete or continuous. A **probability mass function** (PMF) describes the probabilistic behavior of a discrete random variable Y , and is given by

$$f(y) = P(Y = y)$$

for all $y \in A$. Thus, for a given value of y , the PMF is the probability that Y takes on that particular value. As such, the PMF allocates probability to every element in the support set, and hence every outcome of the random process for which it is defined. The left plot in Figure A.1 is a visual display of the probability distribution describing the random variable Y defined in Example 1. With this we can calculate things like the probability that exactly 6 individuals respond to the survey ($P(Y = 6)$), or the probability that 10 or more individuals respond to the survey ($P(Y \geq 10)$). By summing the heights of the bars corresponding to all values of y consistent with these events, we find that $P(Y = 6) = 0.1795$ and $P(Y \geq 10) = 0.0611$. These calculations are depicted visually in the left and right panels of Figure A.2.

A **probability density function** (PDF) describes the probabilistic behavior of a continuous random variable Y . Unlike the probability mass function, which for a particular value of y is itself a probability, we think of the PDF $f(y)$ as being the equation of a **density curve** and probabilities concerning Y are calculated as areas beneath this curve. For instance, a hypothetical probability density function describing the lifetime of an iPhone battery (as in Example 2) is plotted in the right panel of Figure A.1. If we are interested in the probability that an iPhone battery will last up to 10 hours ($P(Y \leq 10)$) or more than 20 hours ($P(Y > 20)$), we calculate the area beneath the curve to the left of 10 and right of 20, respectively. Mathematically this requires integration of the PDF. The two probabilities




Figure A.2: Left: $P(Y = 6) = f(6)$; Right: $P(Y \geq 10) = \sum_{y=10}^{30} f(y)$

of interest in this case are given by 0.8647 and 0.0183 and visualized in the left and right panels of Figure A.3, respectively.




Figure A.3: Left: $P(Y \leq 10) = \int_0^{10} f(y)dy$; Right: $P(Y > 20) = \int_{20}^{\infty} f(y)dy$

While a probability distribution is most efficiently summarized by a plot, such as those given in Figure A.1, the probability function (and hence the distribution) may also be characterized by a closed-form expression. This is the case for several well-known probability distributions which are useful for describing a host of real-life random phenomenon. We dis-

cuss some of these distributions here, focusing on ones that are used routinely in the context of experimentation.

A.1.2 Relevant Distributions

The Binomial Distribution: As noted above, discrete distributions typically describe the randomness associated with counting events. The binomial distribution is one such distribution, and is relevant when counting events in the context of **Bernoulli trials**. Note that a Bernoulli trial is a random process in which there are just two possible outcomes, arbitrarily labelled *successes* and *failures*. Additionally, the occurrence of these outcomes must be independent of one another (i.e., the outcome of one trial does not influence the outcome of any other trial) and the probability of success π (and hence the probability of failure $1 - \pi$), must be the same on each trial. Flipping a coin is a common example of a Bernoulli trial where, for example, the coin turning up ‘heads’ qualifies as a success and ‘tails’ qualifies as a failure. If the coin is fair, the probability of a success is $\pi = 0.5$ each time and whether the coin turns up ‘heads’ on one toss does not influence the outcome of any other toss.

In a sequence of n independent Bernoulli trials, each having probability of success π , the binomial random variable Y counts the number of successes, and we denote it by $Y \sim \text{BIN}(n, \pi)$. The probability mass function $f(y)$ for this distribution, which describes the probability of observing exactly y successes in a sequence of n Bernoulli trials, is given by

$$f(y) = P(Y = y) = \binom{n}{y} \pi^y (1 - \pi)^{n-y}$$

and is defined for $y = 0, 1, 2, \dots, n$ and $\pi \in [0, 1]$. In practice, we obtain probabilities of interest by substituting particular values of y into this formula.

Note that as a special case, when $n = 1$, the binomial distribution simplifies to what is known as the **Bernoulli distribution** which is commonly used to describe response variables that are recorded on a binary scale, such as whether or not an experimental unit clicked or did not click a certain button, or whether a survey respondent was male or female.

The probability mass function for the Bernoulli distribution is given by

$$f(y) = P(Y = y) = \pi^y(1 - \pi)^{1-y}$$

where $y = 0, 1$ and again $\pi \in [0, 1]$.

The Normal Distribution: The normal distribution is arguably the most important and most useful distribution in all of probability and statistics. The veracity of this bold claim will become evident as we work through the statistical analyses associated with different types of experiments. For now we motivate its utility in a practical way by simply stating that there are a remarkable number of real-life phenomena that can be well-modeled by a normal distribution.

A random variable Y is said to be normally distributed if it takes on values $-\infty < y < \infty$ in accordance with the following probability density function

$$f(y) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(y-\mu)^2}{2\sigma^2}}$$

where $-\infty < \mu < \infty$ and $\sigma > 0$. We denote this random variable as $Y \sim N(\mu, \sigma^2)$ and remark that the shape of this distribution is completely determined by the parameters μ and σ . In particular, the distribution can qualitatively be described as ‘bell-shaped and symmetrical’ where μ determines the location of the axis of symmetry and σ determines the dispersion, or spread, of the distribution. Figure A.4 depicts a variety of normal density curves for various values of μ and σ and demonstrates that no matter the (μ, σ) combination, the distribution is always centered at μ and its dispersion is controlled by σ , with larger values corresponding to increased dispersion and smaller values corresponding to decreased dispersion. We note in passing that due to a constraint which says that the area beneath a density curve must equal 1, wider distributions are necessarily shorter than thinner distributions. This is also visualized in Figure A.4.

Note that an important special case exists when $\mu = 0$ and $\sigma = 1$; we call the $N(0, 1)$ distribution the **standard normal distribution** and the corresponding random variable is typically denoted by the letter Z . It can be shown that the following transformation, which




Figure A.4: A variety of normal density curves based on different values of μ and σ

is known as *standardization*, can convert any normal random variable $Y \sim N(\mu, \sigma^2)$ into a standard normal random variable $Z \sim N(0, 1)$:

$$Z = \frac{Y - \mu}{\sigma}$$

We will find the standard normal distribution very useful in the context of hypothesis testing.

The Student's t -Distribution: Another continuous distribution that is very useful in the context of hypothesis testing is the t -distribution, sometimes referred to as the “Student’s” t -distribution (after the pseudonym⁴ of William Gosset, the statistician who first derived

⁴Historical Note: William Gosset was an English statistician who worked at the Guinness Brewery in Dublin Ireland in the early 1900’s. Due to a publication ban imposed by Guinness at the time (because of a previous leak of trade secrets), Gosset was forced to publish under the pseudonym *Student*.

it). Like the normal distribution, the t -distribution is ‘bell-shaped and symmetrical’, but unlike the normal distribution the t -distribution is always centered at 0 and its dispersion is determined by a parameter ν called the **degrees of freedom**. A random variable Y that follows a t -distribution with ν degrees of freedom is denoted $Y \sim t_{(\nu)}$ and the corresponding probability density function is given by

$$f(y) = \frac{\Gamma(\frac{\nu+1}{2})}{\sqrt{\nu\pi}\Gamma(\frac{\nu}{2})} \left(1 + \frac{y^2}{\nu}\right)^{-\frac{\nu+1}{2}}$$

for $-\infty < y < \infty$ and ν is a positive integer. Note that $\Gamma(a)$ is referred to as the “gamma function” and is evaluated as

$$\Gamma(a) = \int_0^{\infty} x^{a-1} e^{-x} dx$$

which, if a is a positive integer, is $\Gamma(a) = (a-1)!$.

Figure A.5 depicts various t -distribution density curves and illustrates how dispersion depends on the degrees of freedom. Notably, as the number of degrees of freedom tends to infinity ($\nu \rightarrow \infty$), the t -distribution converges to the black curve. Although outside the scope of this Appendix, it can be shown that this black curve is the standard normal density curve. In other words

$$\lim_{\nu \rightarrow \infty} t_{(\nu)} = N(0, 1)$$

This will become a practically useful result in the context of various hypothesis tests when we are dealing with very large sample sizes, n .

The Chi-Squared Distribution: The chi-squared distribution (also called the χ^2 -distribution) is another continuous distribution useful in the context of hypothesis testing whose shape is dependent upon a parameter ν called the degrees of freedom. A random variable Y that follows a chi-squared distribution with ν degrees of freedom is denoted $Y \sim \chi_{(\nu)}^2$, and its probability density function is given by

$$f(y) = \frac{y^{\frac{\nu}{2}-1} e^{-y/2}}{2^{\nu/2} \Gamma(\frac{\nu}{2})}$$




Figure A.5: A variety of t -distribution density curves based on different numbers of degrees of freedom ν

for $y \geq 0$ and where ν is a positive integer. Figure A.6 depicts a variety of chi-squared density curves corresponding to different values of ν . As we can see, the shape of chi-squared distribution tends to be right-skewed, with a few special cases exhibiting exponential decay.

The F -Distribution: The F -distribution (also called Snedecor's F -distribution, after Ronald A. Fisher and George W. Snedecor) is another continuous distribution useful in the context of hypothesis testing whose shape is dependent upon two parameters ν_1 and ν_2 called the degrees of freedom. A random variable Y that follows an F -distribution with ν_1 and ν_2 degrees of freedom is denoted $Y \sim F(\nu_1, \nu_2)$, and its probability density function is




Figure A.6: A variety of χ^2 -distribution density curves based on different numbers of degrees of freedom ν

given by

$$f(y) = \frac{\Gamma(\frac{\nu_1+\nu_2}{2})}{\Gamma(\frac{\nu_1}{2})\Gamma(\frac{\nu_2}{2})} \left(\frac{\nu_1}{\nu_2}\right)^{\frac{\nu_1}{2}} y^{\frac{\nu_1}{2}-1} \left(1 + \frac{\nu_1}{\nu_2}y\right)^{-\frac{\nu_1+\nu_2}{2}}$$

for $y \geq 0$ and where ν_1 and ν_2 are positive integers. Figure A.7 depicts a variety of F density curves corresponding to the different values of ν_1 and ν_2 . As we can see, like the chi-squared distribution, the shape of the F -distribution tends to be right-skewed, with a few special cases exhibiting exponential decay.




Figure A.7: A variety of F -distribution density curves based on different numbers of degrees of freedom ν_1 and ν_2

A.1.3 Expectation and Variance

Figures A.4, A.5, A.6, and A.7 demonstrate the variety of different shapes that a probability distribution can exhibit. Not only are these images visually pleasing, they are informative; with one glimpse we can tell which values of y seem typical and which seem extreme, we get sense of how dispersed the distribution is, and we can tell whether it is symmetrical or skewed. However, these observations – when gleaned from a plot – are informal. A quantitative method of communicating the shape of a distribution is with its **moments**. Before discussing moments, however, we must discuss the notion of **expectation**.

The **expected value** of a random variable Y , denoted $E[Y]$, is thought of as the ‘average’

value of Y and as a measure of center in Y 's distribution. Mathematically, the expected value of Y is calculated as

$$E[Y] = \sum_{all\ y} yf(y)$$

if Y is a discrete random variable and as

$$E[Y] = \int_{all\ y} yf(y)dy$$

if Y is a continuous random variable.

Moments, then, are defined to be special expected values, which when taken together, completely specify the shape of a distribution. We define the k^{th} moment of Y to be $E[Y^k]$, which is calculated as in the preceding equations except that y^k (and not y) is multiplied by $f(y)$. Of particular importance in probability and statistics are the first four moments:

- The **first moment** $E[Y]$ quantifies the center of the distribution of Y
- The **second moment** $E[Y^2]$ quantifies the spread of the distribution of Y
- The **third moment** $E[Y^3]$ quantifies the skewness of the distribution of Y
- The **fourth moment** $E[Y^4]$ quantifies the kurtosis (or ‘tailedness’) of the distribution of Y

These four moments provide a tremendous amount of information about the distribution of Y . That said, in practice the first two moments are the ones used most frequently to describe a distribution's shape; relatively speaking more readily useful information is contained in the first two moments than in the others.

While the second moment $E[Y^2]$ itself provides information about the dispersion of a distribution, it is most commonly used in the calculation of the **variance** of Y , $Var[Y]$. The variance of a random variable Y is defined to be

$$Var[Y] = E[(Y - E[Y])^2]$$

Table A.1: Expected values and variances associated with some common distributions

Distribution	$E[Y]$	$Var[Y]$
$Y \sim BIN(n, \pi)$	$n\pi$	$n\pi(1 - \pi)$
$Y \sim N(\mu, \sigma^2)$	μ	σ^2
$Y \sim t_{(\nu)}$	0	$\nu/(\nu - 2)$
$Y \sim \chi^2_{(\nu)}$	ν	2ν
$Y \sim F(\nu_1, \nu_2)$	$\frac{\nu_2}{\nu_2 - 2}$	$\frac{2\nu_2^2(\nu_1 + \nu_2 - 2)}{\nu_1(\nu_2 - 2)^2(\nu_2 - 4)}$

and is interpreted as the expected squared deviation from the mean, with larger values indicating more dispersion and smaller values indicating less dispersion. It can be shown that the equation above can equivalently be expressed as

$$Var[Y] = E[Y^2] - E[Y]^2$$

which makes explicit the dependence of $Var[Y]$ on $E[Y^2]$. Note that the dispersion of a distribution is also commonly communicated in terms of the standard deviation of Y , denoted $SD[Y]$, and calculated as $SD[Y] = \sqrt{Var[Y]}$. Note that Table A.1 contains the expected values and variances of the five distributions described in the previous subsection. As can be seen, these rely entirely on the parameters associated with each distribution.

A.2 Statistical Inference

In practice, we often wish to study a particular characteristic, such as a response variable Y , in some **population** and make inferences about it. In most cases the population is too large to examine in its entirety and so we take a **sample** $\{y_1, y_2, \dots, y_n\}$ from this population and generalize the conclusions drawn in the sample, applying them to the broader population. This process of generalizing sample information to the population from which it was taken is referred to as **statistical inference**. From a probabilistic point of view, we use probability distributions to model sample data, and assume that the chosen distribution is an accurate representation of Y at the population-level.

In the previous section we saw that much of a distribution's information is contained in its shape, and the shape of a given distribution relies entirely on one or more parameters. For

instance, the binomial distribution depends on π , the normal distribution depends on μ and σ , both the t -distribution and the chi-squared distribution rely on degrees of freedom ν , and the F -distribution relies on two types of degrees of freedom, ν_1 and ν_2 . In practice, however, the values of these parameters are unknown and interest typically lies in (i) estimating these parameters in light of the observed data, and/or (ii) testing hypotheses about the parameters. Here we discuss both types of statistical inference, but because the analysis of experiments typically involves testing one or more hypotheses of interest, we place more emphasis on (ii).

A.2.1 A Primer on Point and Interval Estimation

When a data scientist says that they are “fitting” a model to some data, what they really mean is:

- They’ve assumed a certain model or probability distribution is appropriate for describing some characteristic or relationship in a population.
- They have collected data (i.e., a sample from the population) with which they intend to study this characteristic or relationship.
- They intend to use the observed data to estimate the unknown parameters associated with the model or distribution.

Thus, the goal of **point estimation** is to use observed data to obtain reasonable values of a model’s unknown parameters (call them θ) that are consistent with the data that were actually observed. Whereas we typically use Greek letters to denote unknown parameters we use Greek letters over scored by a circumflex (a “hat”⁵), i.e., $\hat{\theta}$, to denote its corresponding estimate. In general, a variety of estimation methods may be used to obtain parameter estimates: the method of moments, maximum likelihood estimation and least squares estimation, to name a few. All estimation procedures have advantages and disadvantages, and so it is important to choose the one that is appropriate for your data and your problem.

⁵The notation $\hat{\theta}$ is read “ θ -hat”.

It is also important to distinguish between point estimation and **interval estimation**. In the context of point estimation we use our data to obtain a single estimate of θ . However, if we were to draw a second sample and repeat the exact same estimation procedure we would very likely obtain a slightly different value of $\hat{\theta}$ than before, simply due to sampling variation. Given this sampling variation, how would you know if your estimate is a good one? In other words, how do you know if your estimate is anywhere close to the true, unknown, value of θ ? The reality is that we can't know this. However, rather than calculating just a point estimate of θ , we can also calculate an interval estimate, more commonly known as a **confidence interval**, for θ . Doing so acknowledges that a point estimate, although likely close to the parameter's true value, is probably not exactly equal to the parameter's true value. Such an interval provides a range within which we are reasonably certain the true value of θ lies. Thus in addition to reporting point estimates of a parameter θ it is most informative to also report a confidence interval for θ as well. For a thorough, but introductory, overview of point and interval estimation techniques see [Bain and Engelhardt \(1992\)](#).

A.2.2 A Primer on Hypothesis Testing

In the context of point and interval estimation we treat the parameter θ as completely unknown and something we need to estimate. However, in some circumstances we may have a belief about the value of θ , and we may wish to use sample data to evaluate whether or not that belief seems reasonable. Statistically speaking such a belief is called a **hypothesis** and the use of data to evaluate that belief is referred to as **hypothesis testing**.

Suppose we believe $\theta = \theta_0$. A formal hypothesis statement corresponding to this can be framed as

$$H_0: \theta = \theta_0 \text{ vs. } H_A: \theta \neq \theta_0$$

We call H_0 the **null hypothesis** and it is the statement we believe to be true, and that we want to test using observed data. The statement denoted H_A is called the **alternative**

hypothesis and it is the complement of H_0 . Thus, exactly one statement is true – either the null hypothesis or the alternative hypothesis – and we use observed data to try and empirically uncover the truth. Note that according to H_A values of θ both larger and smaller than θ_0 correspond to H_0 being false, and so we call such a test **two-sided**. This is to be contrasted with **one-sided** tests for which values of θ larger than θ_0 *or* values of θ smaller than θ_0 (but not both) correspond to H_0 being false. One-sided hypotheses can be stated as

$$H_0: \theta \leq \theta_0 \text{ vs. } H_A: \theta > \theta_0$$

or

$$H_0: \theta \geq \theta_0 \text{ vs. } H_A: \theta < \theta_0$$

depending on the context of the problem and the question that the hypothesis test is designed to answer. No matter which hypothesis is appropriate, the goal is always the same: based on the observed data, we will decide to *reject* H_0 or *not reject* H_0 .

In order to draw such a conclusion, we define a **test statistic** T which is a random variable that satisfies three properties: (i) it must be a function of the observed data, (ii) it must be a function of the parameter θ , and (iii) its distribution must not depend on θ . Assuming the null hypothesis is true, the test statistic T follows a particular distribution which we call the **null distribution**. We then calculate t , the observed value of the test statistic, by substituting the observed data and the hypothesized value of θ into the expression for T . Note that expressions for t commonly incorporate terms of the form $\hat{\theta} - \theta_0$ or $\hat{\theta}/\theta_0$. and so the data enter the expression through the parameter's estimate $\hat{\theta}$.

Next we evaluate the extremity of t relative to the null distribution. If t seems very extreme, as though it is very unlikely to have come from the null distribution, then this gives us reason to believe that the null distribution may not be appropriate. On the other hand, if t appears as though it could have come from the null distribution, then there is no reason to believe the null distribution is inappropriate. The left and right panels of Figure [A.8](#) illustrate these two cases. On the left, the value of t is not at all unreasonable in the

context of the null distribution. However, on the right, the value of t is very extreme and would have been very unlikely if the null distribution (and hence the null hypothesis) really were true. Thus when we observe very extreme values of a test statistic it provides evidence against the null hypothesis, and leads us to believe that perhaps H_0 is not true; and the more extreme t is, the more evidence we have against H_0 . With enough evidence (i.e., extreme enough t) we will choose to reject the null hypothesis.




Figure A.8: Left: A non-extreme value of a test statistic; Right: An extreme value of a test statistic

We formalize the extremity of t using the **p-value** of the test. Probabilistically speaking, a p-value is defined to be the probability of observing a value of the test statistic *at least as extreme* as the value we observed, if the null hypothesis is true. Thus the p-value formally quantifies how “extreme” the observed test statistic is. Whether large values of t , small values of t , or both, are to be considered extreme depends on whether H_A is one- or two-sided. When H_A is two-sided, both large and small values of t are considered extreme and we define the p-value mathematically as

$$\text{p-value} = P(T \geq |t|) + P(T \leq -|t|)$$

which, if the null distribution is symmetrical, is equivalent to $2P(T > |t|)$. The left panel of Figure A.9 provides a visual depiction of this calculation.

When H_A is one-sided then either large values of t or small values of t are considered extreme, and this depends on the direction of the inequality in H_A . If $H_A: \theta > \theta_0$, values of θ larger than θ_0 and hence large values of t will render H_0 false. Thus in this case large values of t are considered extreme and the p-value is calculated as

$$\text{p-value} = P(T \geq t).$$

The center panel of Figure A.9 provides a visual depiction of this calculation. If $H_A: \theta < \theta_0$, values of θ smaller than θ_0 and hence small values of t will render H_0 false. Thus in this case small values of t are considered extreme and the p-value is calculated as

$$\text{p-value} = P(T \leq t).$$

The right panel of Figure A.9 provides a visual depiction of this calculation.




Figure A.9: Illustration of the calculation of p-values in one- and two-sided tests

How “extreme” t must be, and hence how small the p-value must be to reject H_0 , is determined by the **significance level** of the test, which we denote by α . In particular, if

- p-value $\leq \alpha$ we reject H_0
- p-value $> \alpha$ we do not reject H_0

Note that $\alpha = 0.01$ or 0.05 are common choices. In order to motivate these choices we need

to discuss the two types of error that can be made when drawing conclusions in the context of a hypothesis test.

Recall that by design either H_0 is true or H_A is true. This means that there are four possible outcomes when using data to decide which statement is true:

- (1) H_0 is true and we correctly do not reject it
- (2) H_0 is true and we incorrectly reject it
- (3) H_0 is false and we incorrectly do not reject it
- (4) H_0 is false and we correctly reject it

Obviously scenarios (1) and (4) are ideal since in them we are making the correct decision, and (2) and (3) should be avoided since in those scenarios we are not making the correct decision. Scenarios (2) and (3) are respectively referred to as **Type I error** and **Type II error**. Clearly we would like to reduce the likelihood of making either type of error, but it is important to recognize that in practice there are different consequences to each type of error, and so we may wish to treat them differently. To make this point clear, consider a courtroom analogy where the defendant is assumed innocent until proven guilty. This hypothesis can be stated formally as

$$H_0: \text{the defendant is innocent} \text{ vs. } H_A: \text{the defendant is guilty}$$

Within this analogy a Type I error occurs when the defendant is truly innocent, but the evidence leads the jury to find the defendant guilty. Thus, this error leads to an innocent person being convicted of a crime they did not commit. A Type II error, on the other hand, occurs when the defendant is truly guilty, but the evidence leads the jury to find the defendant innocent. In this case the error leads to a criminal being set free. In this analogy, and in any hypothesis testing setting, both types of errors lead to negative outcomes, but these negative outcomes may be prioritized differently.

Fortunately it is possible to control the frequency with which these types of errors occur. We do so by controlling the significance level and **power** of the test. We define a test's significance level to be $\alpha = P(\text{Type I Error})$ and we define the power of a test to $1 - \beta$ where $\beta = P(\text{Type II Error})$. Thus it is desirable to have a test with a small significance level and a large power since this corresponds to simultaneously reducing both types of errors.

In practice we choose α and β based on how often we are comfortable allowing Type I and Type II errors to occur. For instance, if we can only tolerate a Type I error 1% of the time, then we would choose $\alpha = 0.01$ and if we can only tolerate making a Type II error 5% of the time, then we would choose $\beta = 0.05$. With these choices we would say that the corresponding hypothesis test has a 1% significance level and 95% power. Common choices for significance level and power are respectively 5% and 80%, corresponding to $\alpha = 0.05$ and $\beta = 0.2$.

As is now apparent, the significance level α (i.e., the probability of making a Type I error), determines how small a p-value must be (and hence how extreme t must be) in order to reject a null hypothesis. This decision should be made prior to testing the hypothesis and in fact prior to collecting any data. We defer a discussion of controlling power until Chapter 2 where we will see that for a fixed value of α the power determines the sample size and so it is also a decision that should be made prior to collecting the data, else you will not know how much data to collect.

A.3 Linear Regression

In this section we provide a brief overview of **linear regression**. Linear regression is a form of statistical modeling that is appropriate when interest lies in relating a response variable (Y) to one or more explanatory variables (x_1, x_2, \dots, x_p) . The idea is that Y is influenced in some manner by the explanatory variables through an unknown function $f(\cdot)$:

$$Y = f(x_1, x_2, \dots, x_p).$$

The purpose of statistical modeling in general, and linear regression in particular, is to approximate this function $f(\cdot)$. The typical linear regression model in this situation is given by

$$Y = \beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p + \epsilon$$

where Y is the response variable; the x_j 's are explanatory variables which we treat as fixed (not random) quantities; the β 's are unknown parameters that quantify the influence of a particular explanatory variable on the response; and $\epsilon \sim N(0, \sigma^2)$ is a random error term that accounts for the fact that $f(x_1, x_2, \dots, x_p) \neq \beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p$ exactly. The distributional assumption for ϵ has several consequences. Chief among them is that Y is also a random variable and

$$Y \sim N(\mu = \beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p, \sigma^2).$$

Thus, for particular values of the explanatory variables, we expect the response variable to be equal to $\mu = \beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p$, on average. Variation around this relationship is quantified by σ^2 .

Based on this distributional assumption we find that $E[Y|x_1 = x_2 = \cdots = x_p = 0] = \beta_0$, and so β_0 is interpreted as the intercept of the model – the expected response when all of the explanatory variables are equal to zero. Also, notice that

$$\begin{aligned} E[Y|x_j = x + 1] - E[Y|x_j = x] &= (\beta_0 + \beta_1 x_1 + \cdots + \beta_j(x + 1) + \cdots + \beta_p x_p) \\ &\quad - (\beta_0 + \beta_1 x_1 + \cdots + \beta_j x + \cdots + \beta_p x_p) \\ &= \beta_0 + \beta_1 x_1 + \cdots + \beta_j x + \beta_j + \cdots + \beta_p x_p \\ &\quad - \beta_0 - \beta_1 x_1 - \cdots - \beta_j x - \cdots - \beta_p x_p \\ &= \beta_j \end{aligned}$$

As such, β_j is interpreted as the expected change in response associated with a unit increase in x_j ($j = 1, 2, \dots, p$), while holding all other explanatory variables fixed. Given the intuitively pleasing interpretations of these coefficients, it should be clear that linear regression models are well-suited for *explanatory modeling*, although they may also be used effectively for *predictive modeling*. See [Shmueli \(2010\)](#) for an interesting discussion of these two goals of

statistical modeling.

Whether we wish to use a linear model for explanatory or predictive purposes, we need to estimate the regression coefficients. Recall the β 's are unknown parameters. This is typically done with **least squares estimation** where the goal is to find the values of $\beta_0, \beta_1, \dots, \beta_p$ that minimize the error, ϵ , associated with the model. Specifically, for observed data given by $(y_i, x_{i1}, x_{i2}, \dots, x_{ip})$, $i = 1, 2, \dots, n$, we wish to minimize

$$\sum_{i=1}^n \epsilon_i^2 = \sum_{i=1}^n (y_i - (\beta_0 + \beta_1 x_{i1} + \dots + \beta_p x_{ip}))^2$$

with respect to $\boldsymbol{\beta} = (\beta_0, \beta_1, \dots, \beta_p)^T$.

By writing the linear regression model above in matrix form as

$$\mathbf{y} = X\boldsymbol{\beta} + \boldsymbol{\epsilon},$$

it can be shown that the least squares estimate of $\boldsymbol{\beta}$ and hence the individual β 's is given by $\hat{\boldsymbol{\beta}} = (X^T X)^{-1} X^T \mathbf{y} = (\hat{\beta}_0, \hat{\beta}_1, \dots, \hat{\beta}_p)^T$ where

- $\mathbf{y} = (y_1, y_2, \dots, y_n)^T$ is an $n \times 1$ vector of response variable observations,
- $\boldsymbol{\epsilon} = (\epsilon_1, \epsilon_2, \dots, \epsilon_n)^T$ is an $n \times 1$ vector of random errors,
- $\boldsymbol{\beta} = (\beta_0, \beta_1, \dots, \beta_p)^T$ is a $(p + 1) \times 1$ vector of regression coefficients, and
- X is the following $n \times (p + 1)$ matrix of explanatory variable observations

$$X = \begin{bmatrix} 1 & x_{11} & x_{12} & \cdots & x_{1p} \\ 1 & x_{21} & x_{22} & \cdots & x_{2p} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n1} & x_{n2} & \cdots & x_{np} \end{bmatrix}$$

With the regression coefficients estimated we define the **fitted values** $\hat{\mu}_i = \hat{\beta}_0 + \hat{\beta}_1 x_{i1} + \dots + \hat{\beta}_p x_{ip}$ to be the estimated expected response for specific values of the explanatory

variables. Next we define the **residuals** $e_i = y_i - \hat{\mu}_i$ to be the difference between the observed value of the response and what the model predicts the response to be. It can be shown that the least squares estimate of σ^2 is based on the residuals, and in particular is given by

$$\hat{\sigma}^2 = \frac{\sum_{i=1}^n e_i^2}{n - p - 1} = \frac{\sum_{i=1}^n (y_i - \hat{\mu}_i)^2}{n - p - 1}.$$

This estimate is sometimes referred to as the **mean squared error** (*MSE*) of the model, since $\hat{\sigma}$ quantifies the typical distance (error) between an observed response value and the value predicted by the model.

Having estimated $\beta_0, \beta_1, \dots, \beta_p$ and σ^2 , the fitted linear regression model can be used for inference and prediction. Of particular importance are hypothesis tests for the individual β 's. For instance, the hypothesis $H_0: \beta_j = 0$ vs. $H_A: \beta_j \neq 0$ may be used to formally evaluate whether the explanatory variable x_j significantly influences Y and whether it belongs in the model. Also of importance are confidence and prediction intervals for predicted values of Y . For a much more thorough (yet approachable) treatment of linear regression see [Abraham and Ledolter \(2006\)](#).

A.4 Logistic Regression

As was shown in the previous section, linear regression is an effective method of modeling the relationship between a single response variable (Y), and one or more explanatory variables (x_1, x_2, \dots, x_p). However, ordinary linear regression assumes that the response variable follows a normal distribution (i.e., $Y \sim N(\mu, \sigma^2)$); when the response variable is binary, this assumption is no longer valid. In the context of a binary response, the Bernoulli distribution (i.e., $Y \sim \text{BIN}(1, \pi)$) is a much more appropriate distributional assumption, but ordinary linear regression is no longer appropriate. Consequently, we use **logistic regression** to model the relationship between a binary response variable and one or more explanatory variables.

In the context of a linear regression model, the model is formulated so that the expected

response (given the values of the explanatory variables) is equated to the **linear predictor** $\beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p$. Namely

$$E[Y|x_1, x_2, \dots, x_p] = \mu = \beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p.$$

In the context of logistic regression we would also like to relate the expected response to the linear predictor, but in this case $E[Y] = \pi$ which is a probability that must lie within $[0, 1]$. However, it is unrealistic to impose this constraint on the linear predictor and so equating π and $\beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p$ does not make sense. Instead, we relate the linear predictor to $E[Y] = \pi$ through a monotonic differentiable **link function** that maps $[0, 1]$ to the real numbers. Logistic regression arises when this link function is chosen to be the **logit** function:

$$\log\left(\frac{\pi}{1 - \pi}\right) = \beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p \quad (\text{A.1})$$

Note that different link functions result in different **generalized linear models**. See [McCullagh and Nelder \(1989\)](#) for a thorough and more general treatment of generalized linear models.

Notice that by inverting the link function the expected response (given the values of the explanatory variables) in a logistic regression is

$$E[Y|x_1, x_2, \dots, x_p] = \pi = \frac{e^{\beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p}}{1 + e^{\beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p}}.$$

To interpret β_0 we consider setting each explanatory variable to zero (i.e., $x_1 = x_2 = \cdots = x_p = 0$) in equation (A.1). In doing so we see that β_0 is the **log-odds** that $Y = 1$, or in other words, that e^{β_0} is the **odds** that the response would equal 1 when $x_1 = x_2 = \cdots = x_p = 0$. The interpretation of β_j , for $j = 1, 2, \dots, p$, is seen by considering equation (A.1) for different values of x_j . In particular, let π_x be the value of π when $x_j = x$ and let π_{x+1} be the value of

π when $x_j = x + 1$, and notice that

$$\begin{aligned}
\log\left(\frac{\pi_{x+1}}{1 - \pi_{x+1}}\right) - \log\left(\frac{\pi_x}{1 - \pi_x}\right) &= (\beta_0 + \beta_1 x_1 + \cdots + \beta_j(x + 1) + \cdots + \beta_p x_p) \\
&\quad - (\beta_0 + \beta_1 x_1 + \cdots + \beta_j x + \cdots + \beta_p x_p) \\
&= \beta_0 + \beta_1 x_1 + \cdots + \beta_j x + \beta_j + \cdots + \beta_p x_p \\
&\quad - \beta_0 - \beta_1 x_1 - \cdots - \beta_j x - \cdots - \beta_p x_p \\
&= \beta_j
\end{aligned}$$

and simplifying the left hand side yields

$$\log\left(\frac{\pi_{x+1}}{1 - \pi_{x+1}} \bigg/ \frac{\pi_x}{1 - \pi_x}\right) = \beta_j$$

which makes it clear that β_j is interpreted as a **log-odds ratio** and hence e^{β_j} is interpreted as the **odds ratio**, comparing the odds that $Y = 1$ when $x_j = x + 1$ vs. when $x_j = x$ (all else being equal). We note in passing that **maximum likelihood estimation** is typically used to estimate the β 's and hence π . This derivation is outside the scope of these notes, so the interested reader is referred to [McCullagh and Nelder \(1989\)](#) for more details.

Just as in the case of linear regression we can perform hypothesis tests to determine whether subsets of the regression coefficients are equal to zero. Practically this amounts to fitting and comparing nested models (full models and reduced models that assume H_0 is true) to determine whether they differ significantly. Because maximum likelihood estimation is used to fit such models, this comparison is made using **likelihood ratio tests**. In order to test individual hypotheses about particular regression coefficients, like $H_0 : \beta_j = 0$, we can similarly apply a likelihood ratio test that compares the models with and without β_j . Alternatively, we may also use individual Z -tests for this purpose. We discuss these tests in more detail in [Section 5.3.2](#).

References

- Abraham, B. and J. Ledolter (2006). *Introduction to regression modeling*. Thomson Brooks/Cole.
- Bain, L. J. and M. Engelhardt (1992). *Introduction to probability and mathematical statistics* (2nd ed.). Brooks/Cole.
- Dunnett, C. W. (1955). A multiple comparison procedure for comparing several treatments with a control. *Journal of the American Statistical Association* 50(272), 1096–1121.
- Edgington, E. and P. Onghena (2007). *Randomization tests*. CRC Press.
- Kuleshov, V. and D. Precup (2000). Algorithms for the multi-armed bandit problem. *Journal of Machine Learning Research* 1, 1–48.
- McCullagh, P. and J. A. Nelder (1989). Generalized linear models, no. 37 in monograph on statistics and applied probability.
- Montgomery, D. C. (2017). *Design and analysis of experiments* (9th ed.). John Wiley & Sons.
- Rosenbaum, P. R. and D. B. Rubin (1983). The central role of the propensity score in observational studies for causal effects. *Biometrika* 70(1), 41–55.
- Scott, S. L. (2010). A modern bayesian look at the multi-armed bandit. *Applied Stochastic Models in Business and Industry* 26(6), 639–658.
- Shapiro, S. S. and M. B. Wilk (1965). An analysis of variance test for normality (complete samples). *Biometrika* 52(3/4), 591–611.

- Shmueli, G. (2010). To explain or to predict? *Statistical Science* 25(3), 289–310.
- Shmueli, G. (2017). Analyzing behavioral big data: Methodological, practical, ethical, and moral issues. *Quality Engineering* 29(1), 57–74.
- Siroker, D. and P. Koomen (2013). *A/B testing: The most powerful way to turn clicks into customers*. John Wiley & Sons.
- Steiner, S. H. and R. J. MacKay (2005). *Statistical Engineering: an algorithm for reducing variation in manufacturing processes*. ASQ Quality Press.
- Welch, B. L. (1947). The generalization of student’s problem when several different population variances are involved. *Biometrika* 34(1/2), 28–35.
- Whittle, P. (1979). Discussion on: bandit processes and dynamic allocation indices [jr statist. soc. b, 41, 148–177, 1979] by jc gittins. *J. Roy. Statist. Soc. Ser. B* 41, 165.